Quiz, week 11

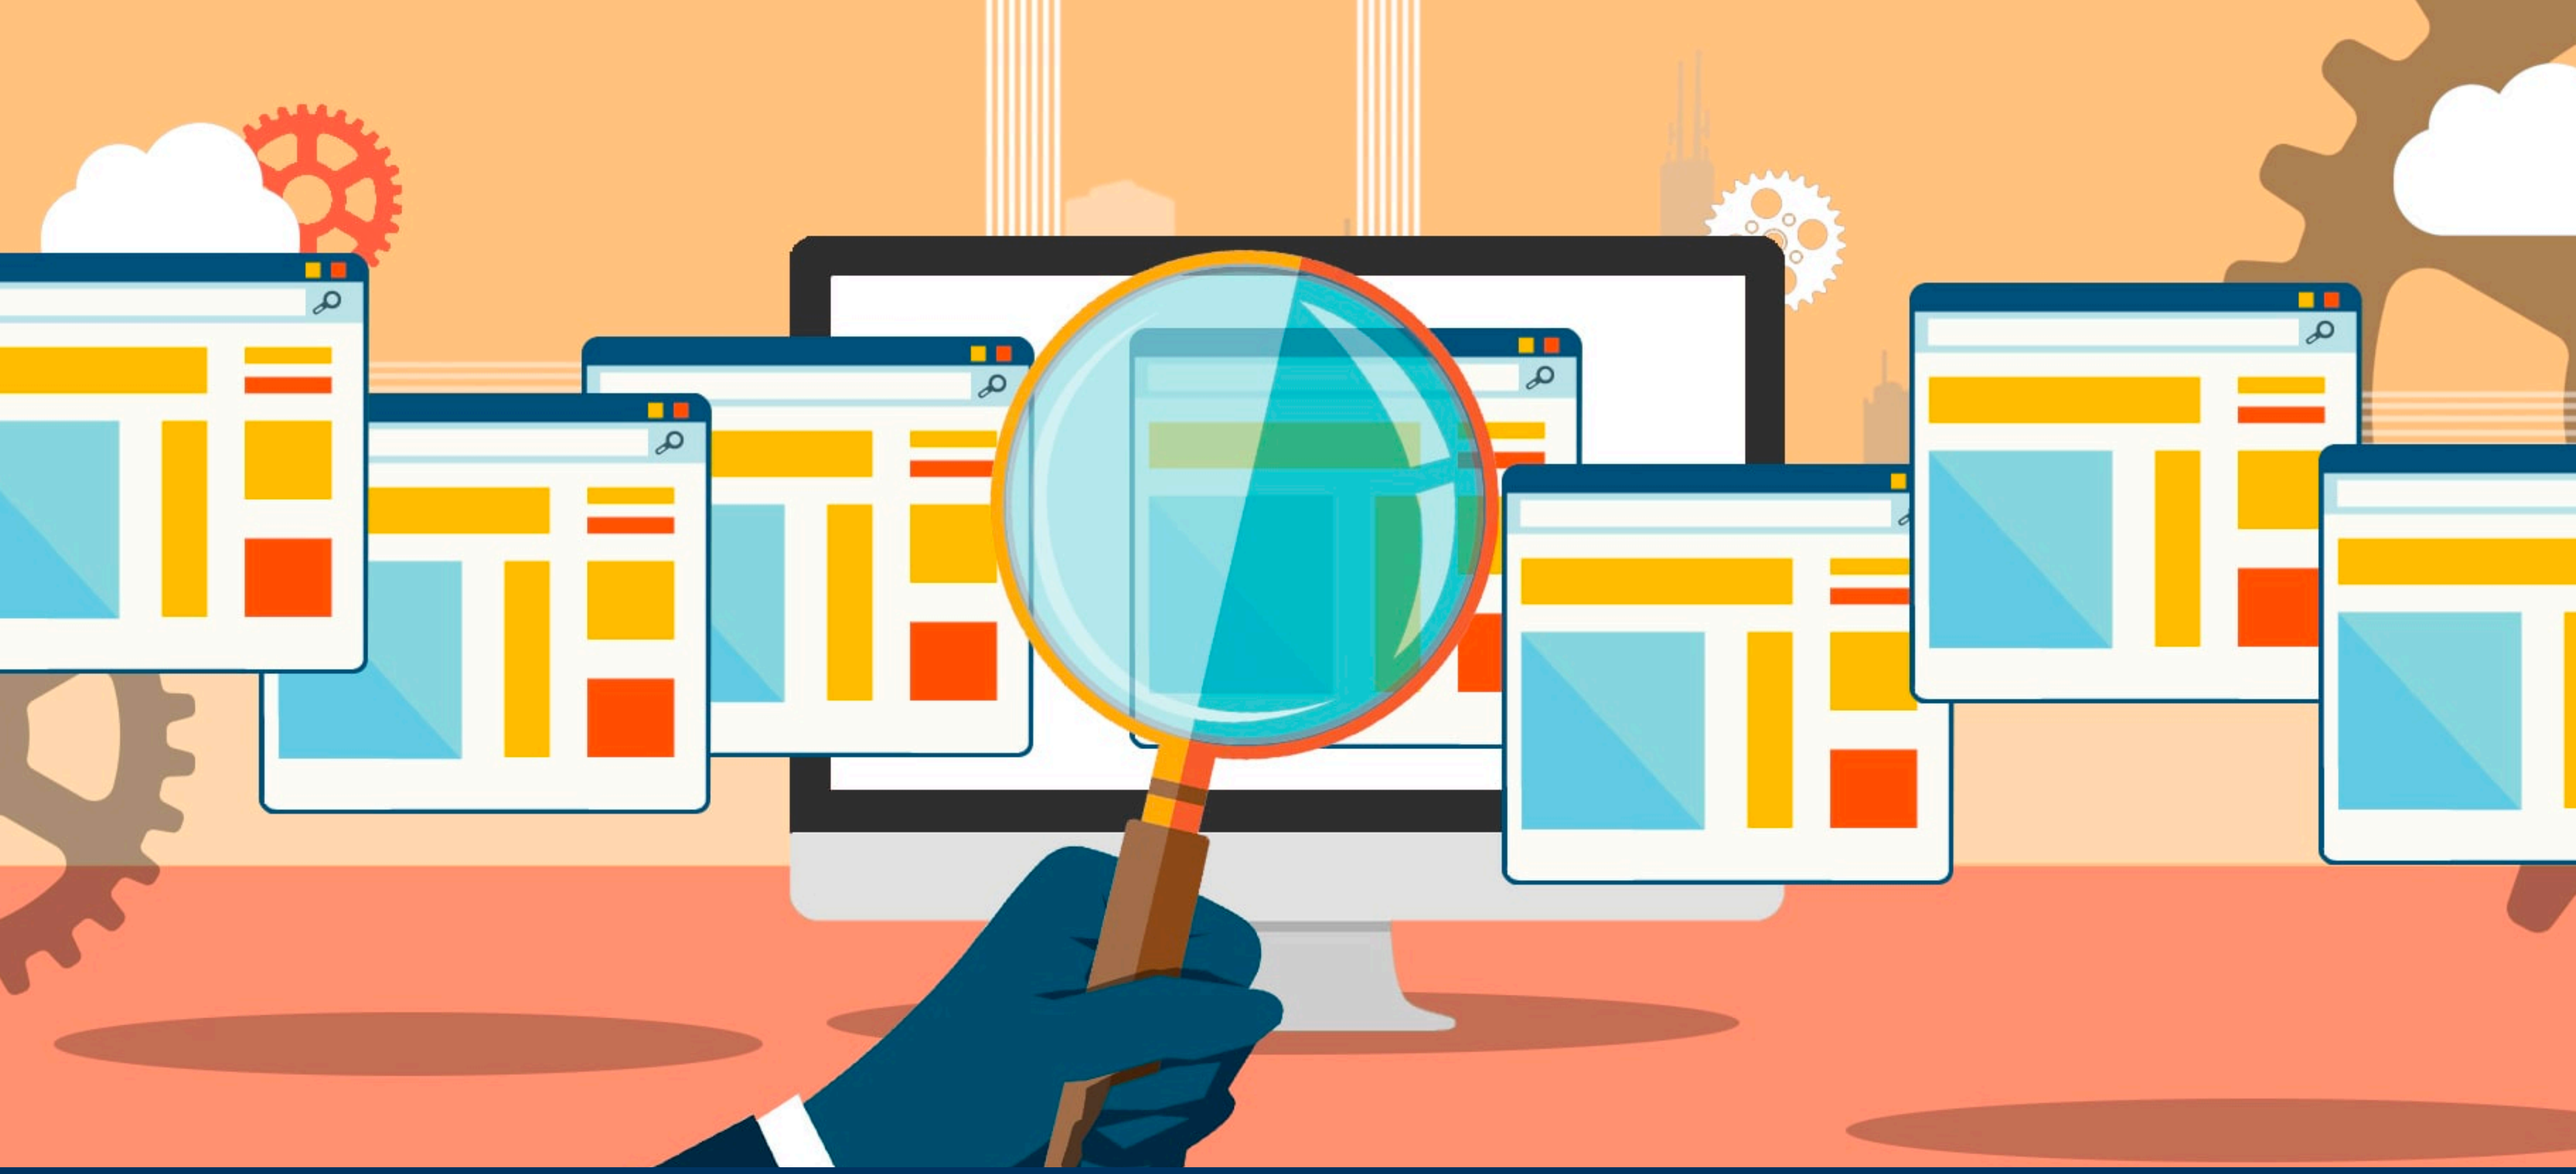

Supervised Learning, pt 2

INST414 - Data Science Techniques

This Module's Learning Objectives

Part 1

- Differentiate between classification and regression in supervised learning
- Describe how voting is used in for k-nearest neighbors classification
- Differentiate binary, multi-class, and multi-label classification
- Define overfitting and describe its impact on generalizability

This Module's Learning Objectives

Part 1

- Differentiate between classification and regression in supervised learning
- Describe how voting is used in for k-nearest neighbors classification
- Differentiate binary, multi-class, and multi-label classification
- Define overfitting and describe its impact on generalizability

Also called the “training data”

What’s a simple rule for classifying new data based on our sample set?

Use the label of the “closest” sample

The “nearest neighbor” approach

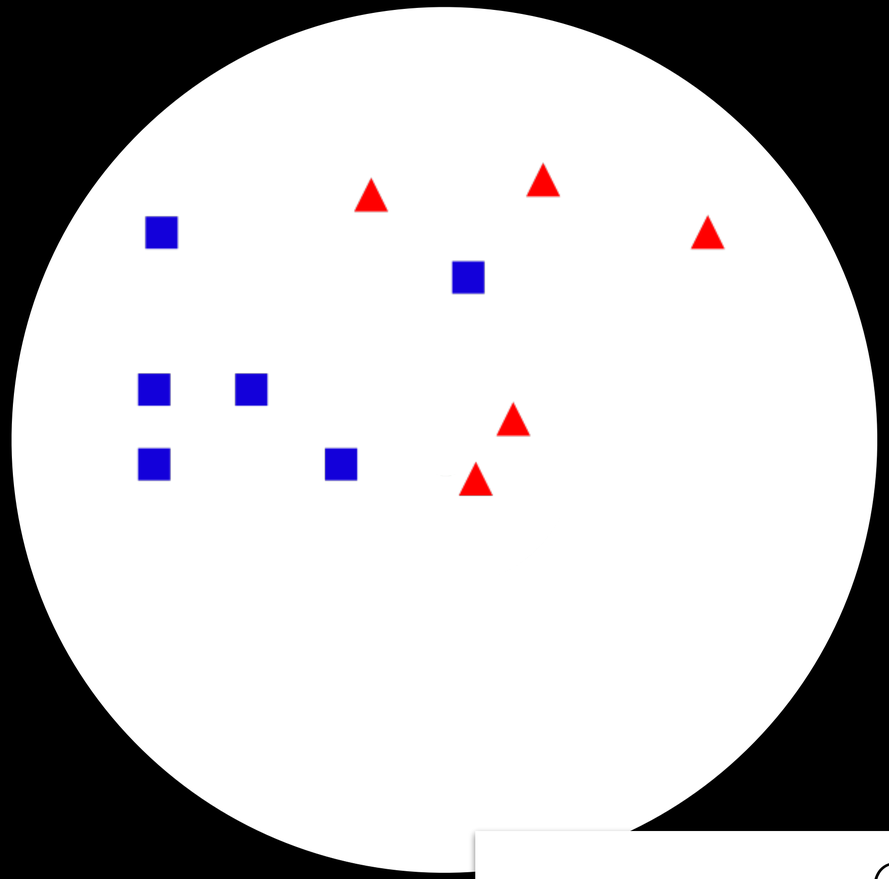

Given a dataset of two classes
(blue/red labels)

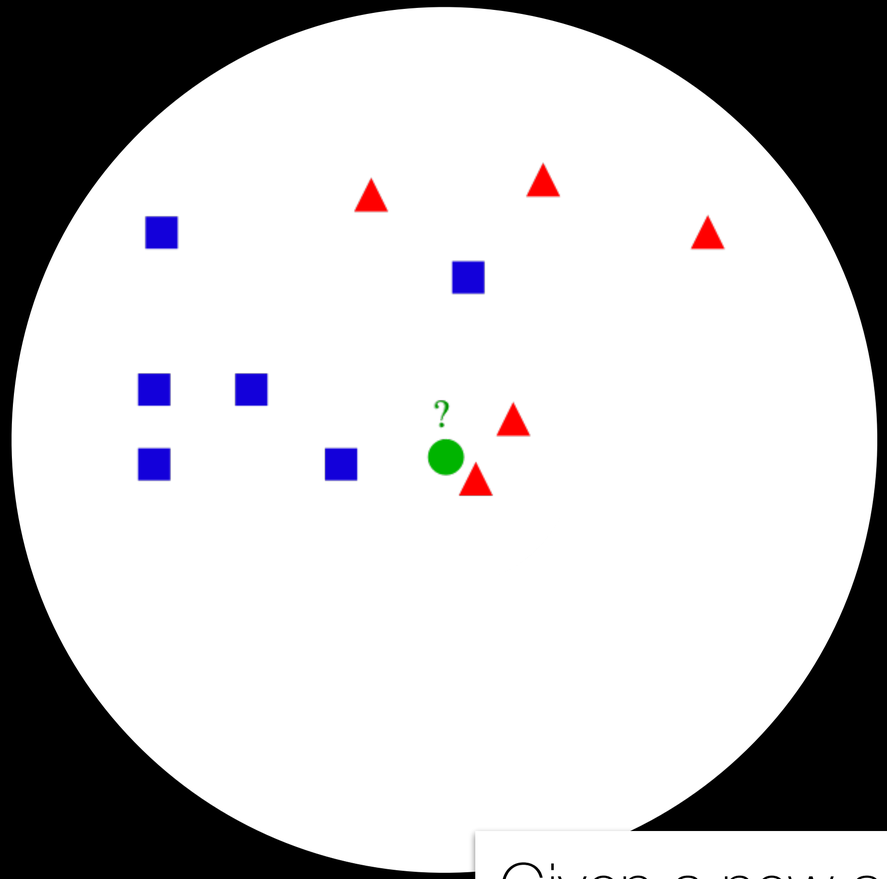

Given a new sample (green data point), what label should it be?

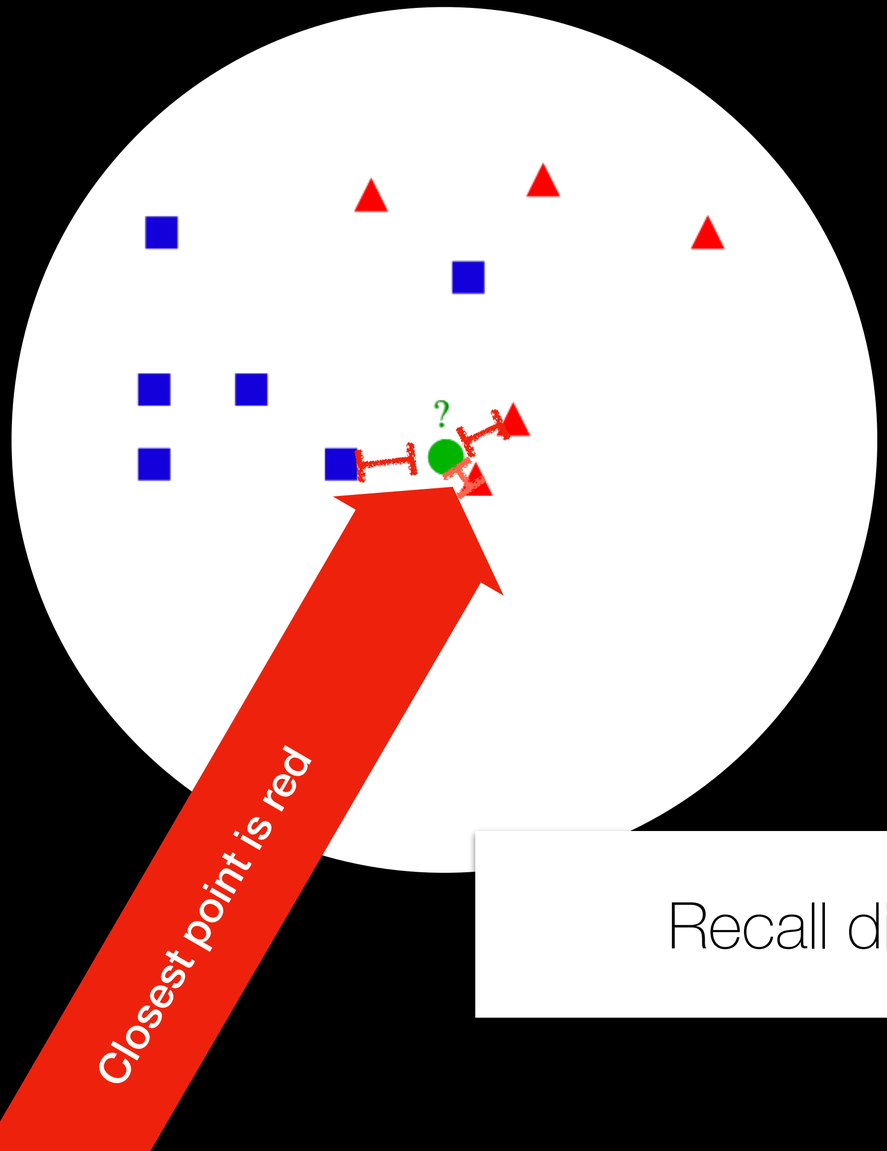

Recall distance metrics from Module 3

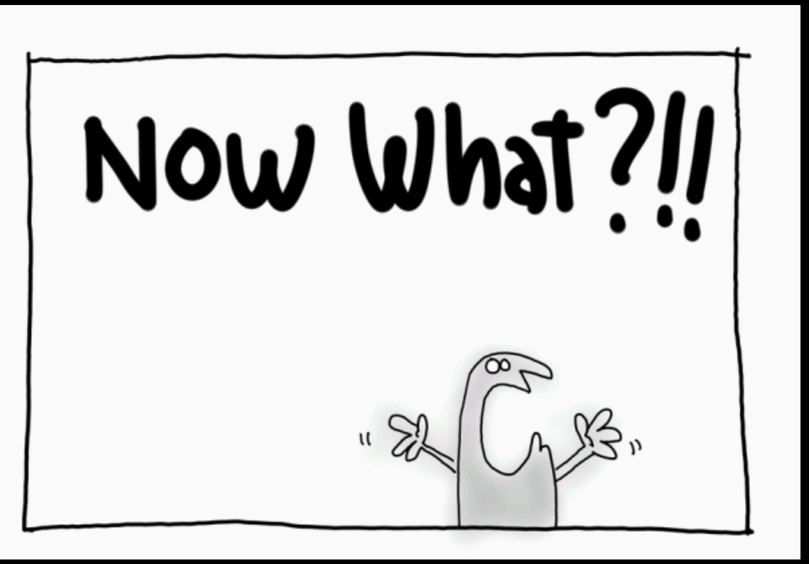

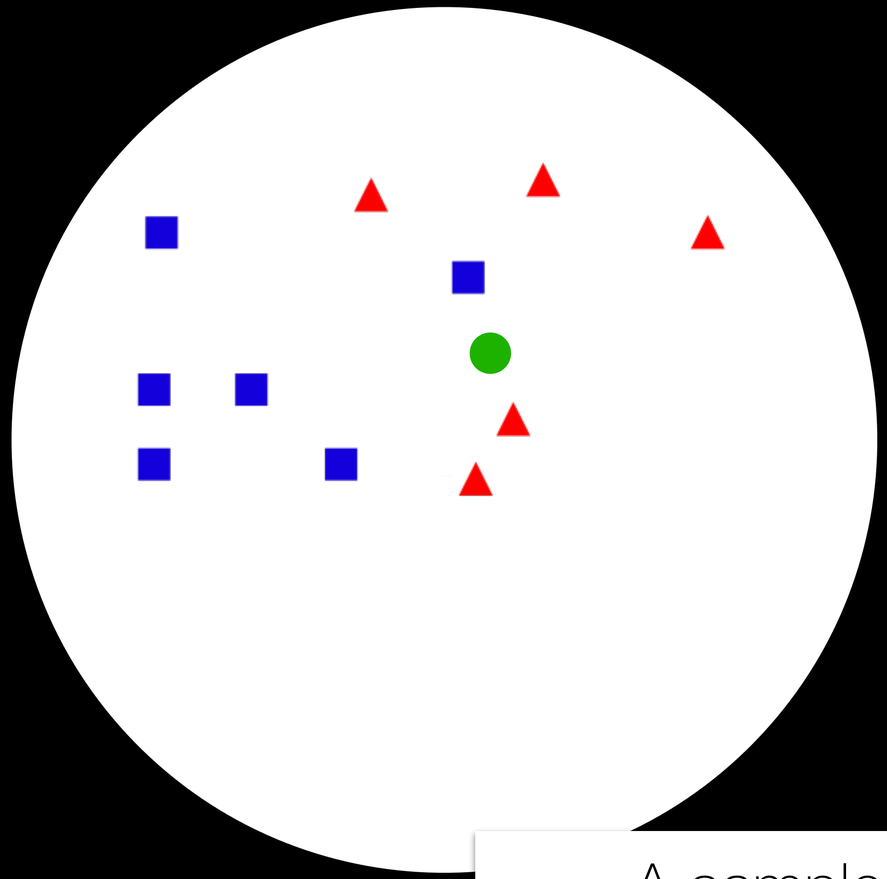

A sample could be equidistant between points

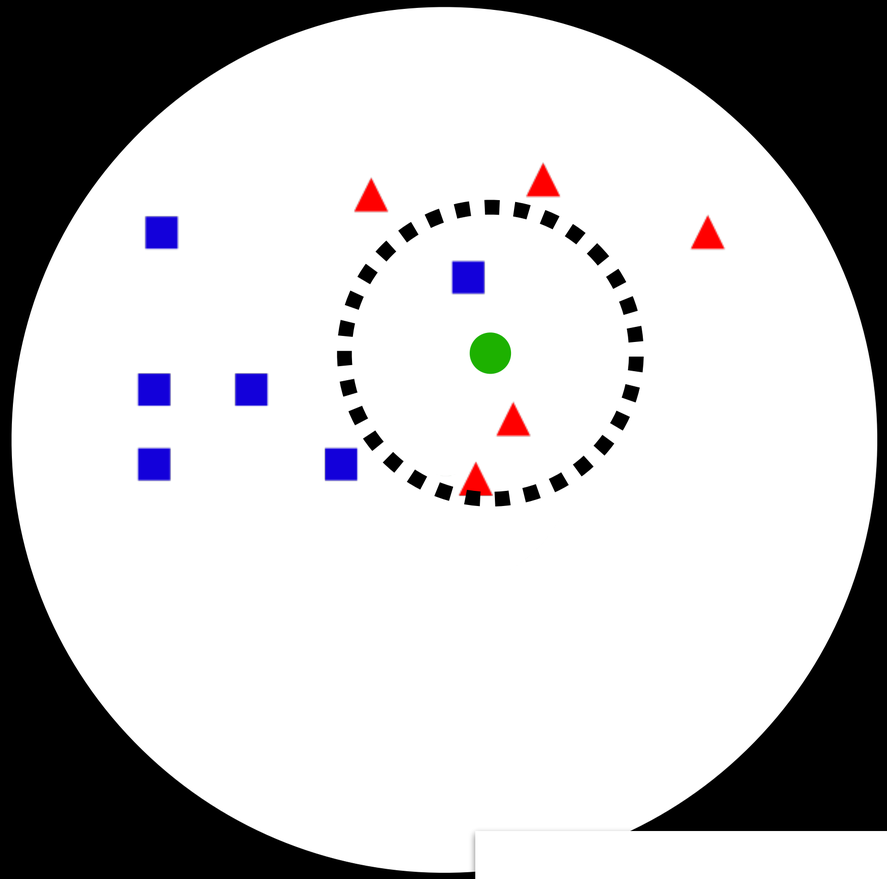

Take a vote among k-nearest neighbors

k=3 —> Red label

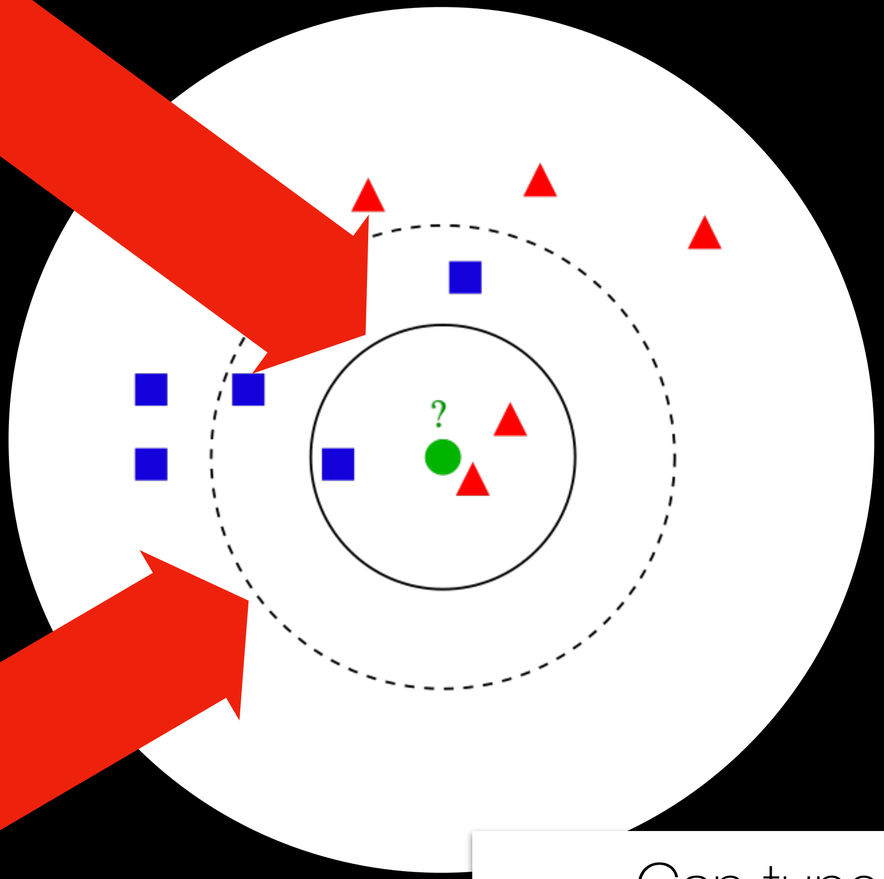

k=5 —> Blue label

Can tune k-value as needed to improve classification accuracy

KNN for Classification (example: 2 features)

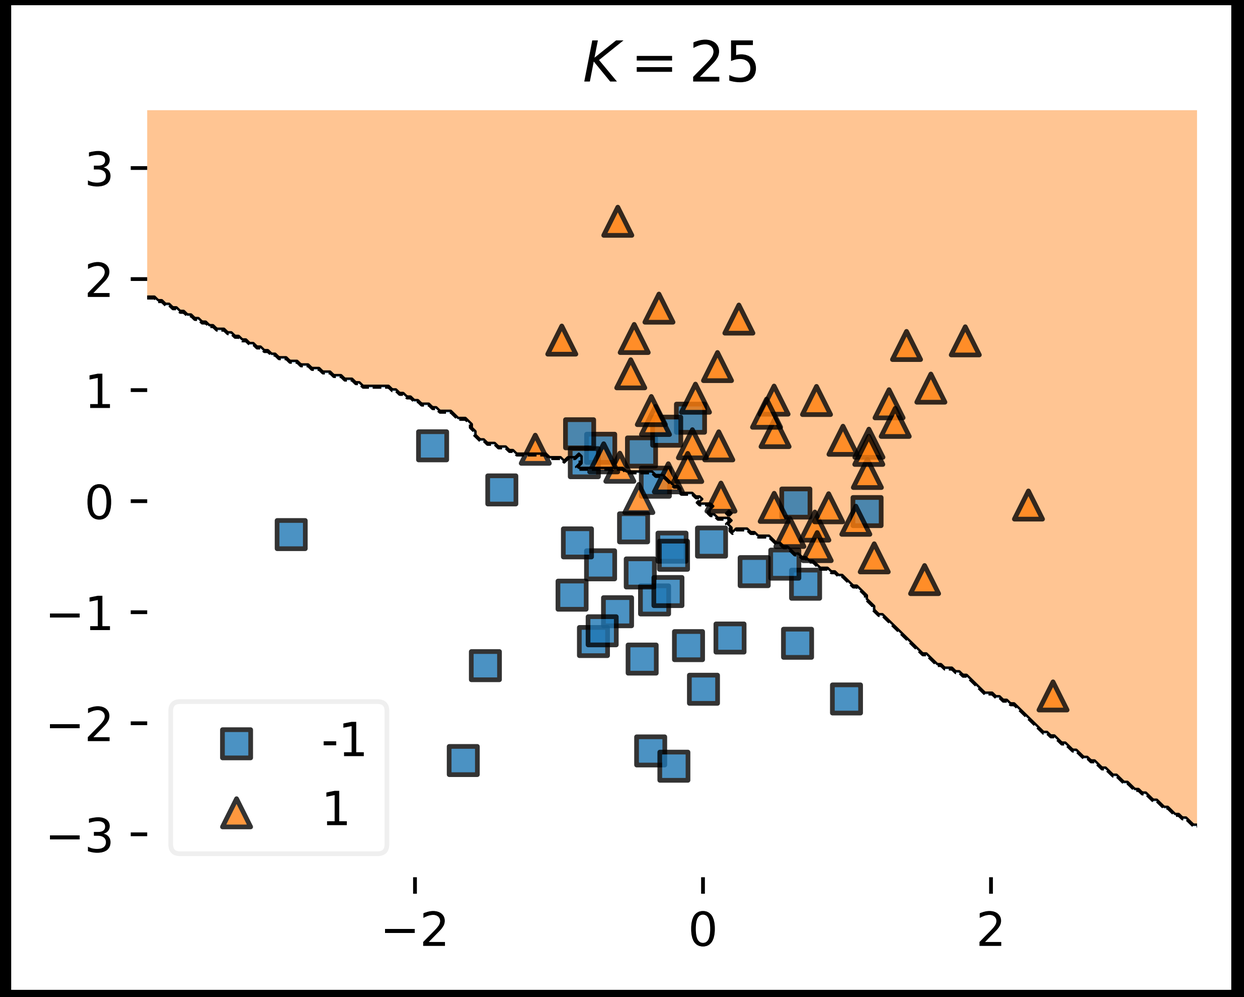

KNN for Classification (example: 2 features)

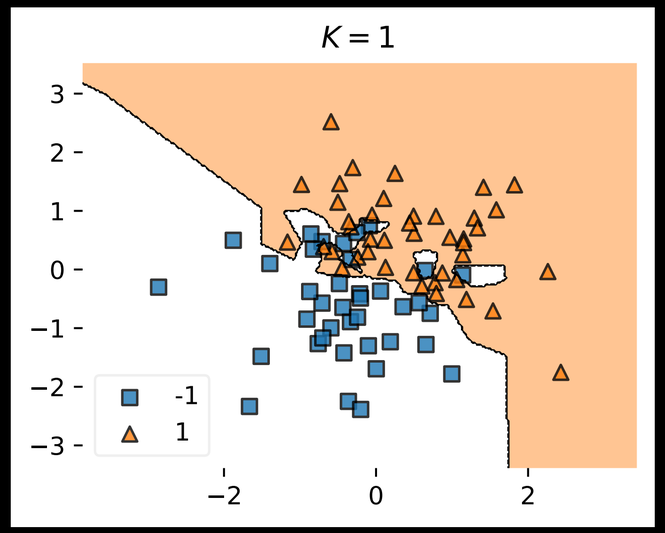

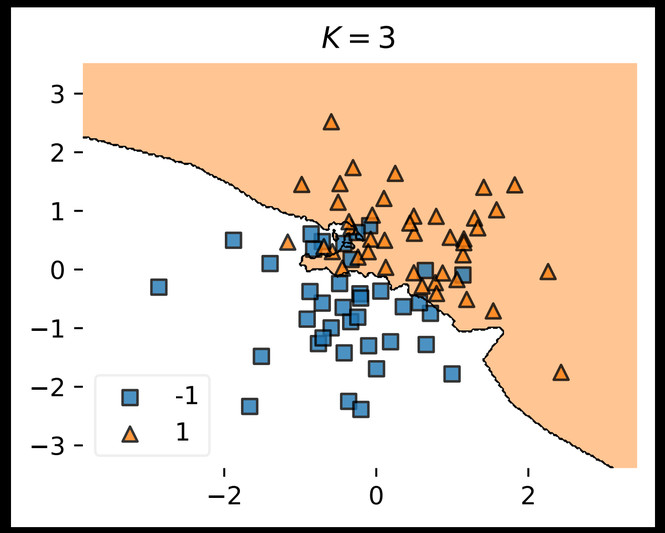

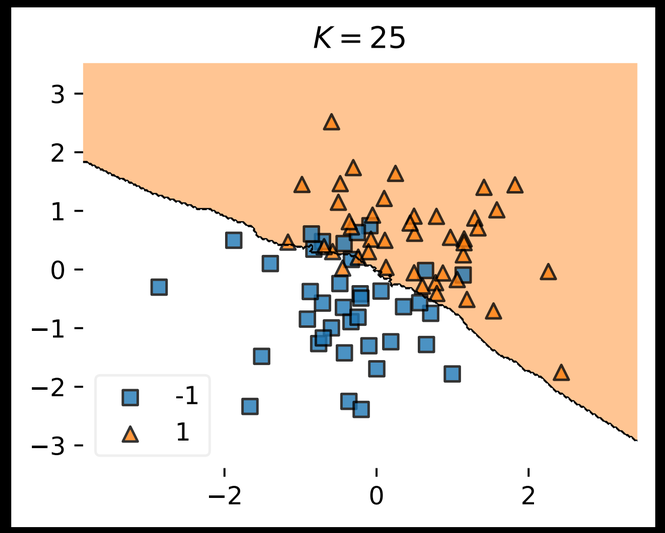

Decision Boundary

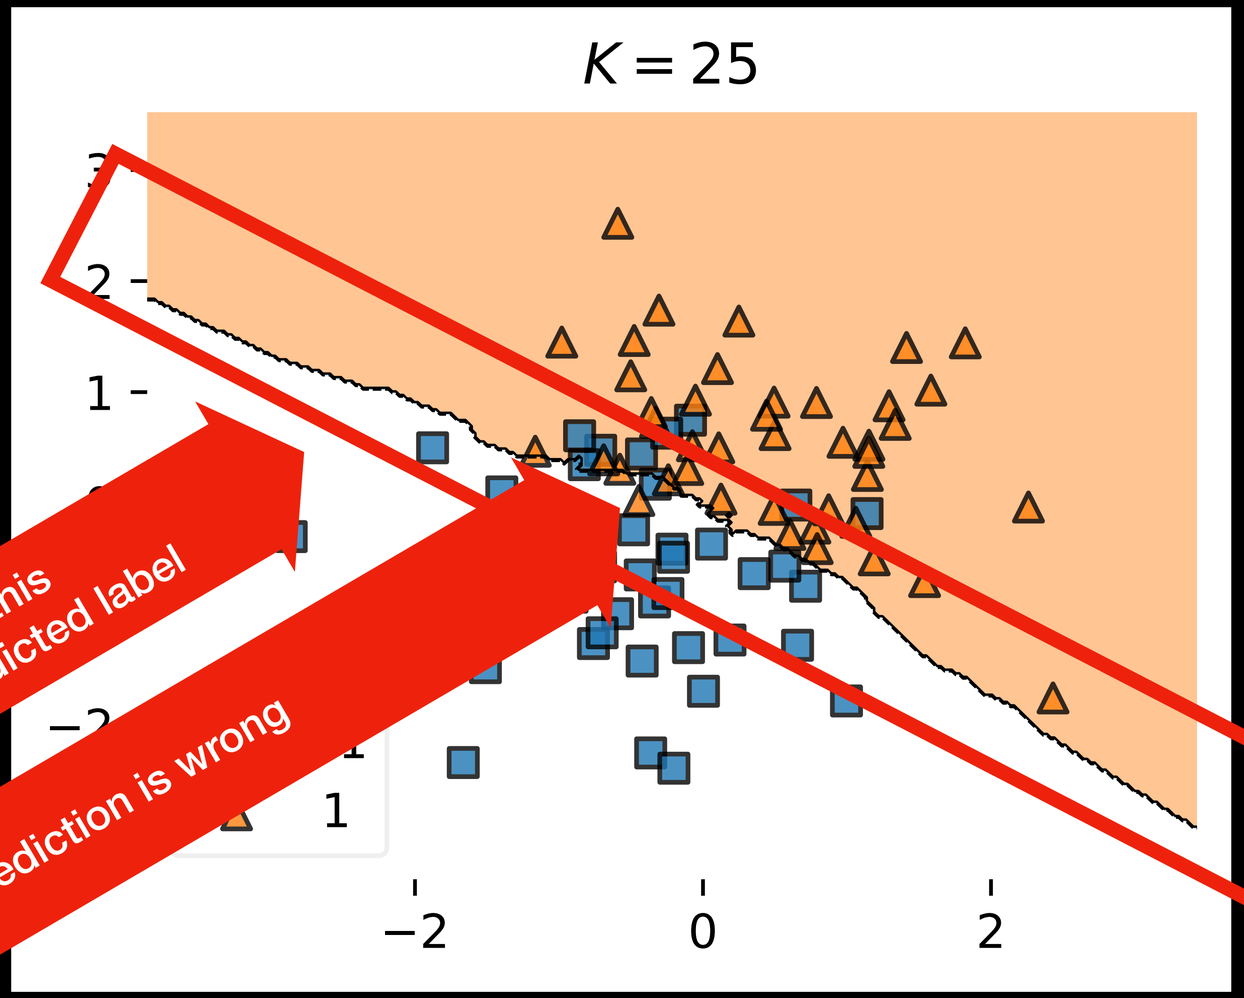

All elements on the same side of this decision boundary have the same predicted label

Event if prediction is wrong

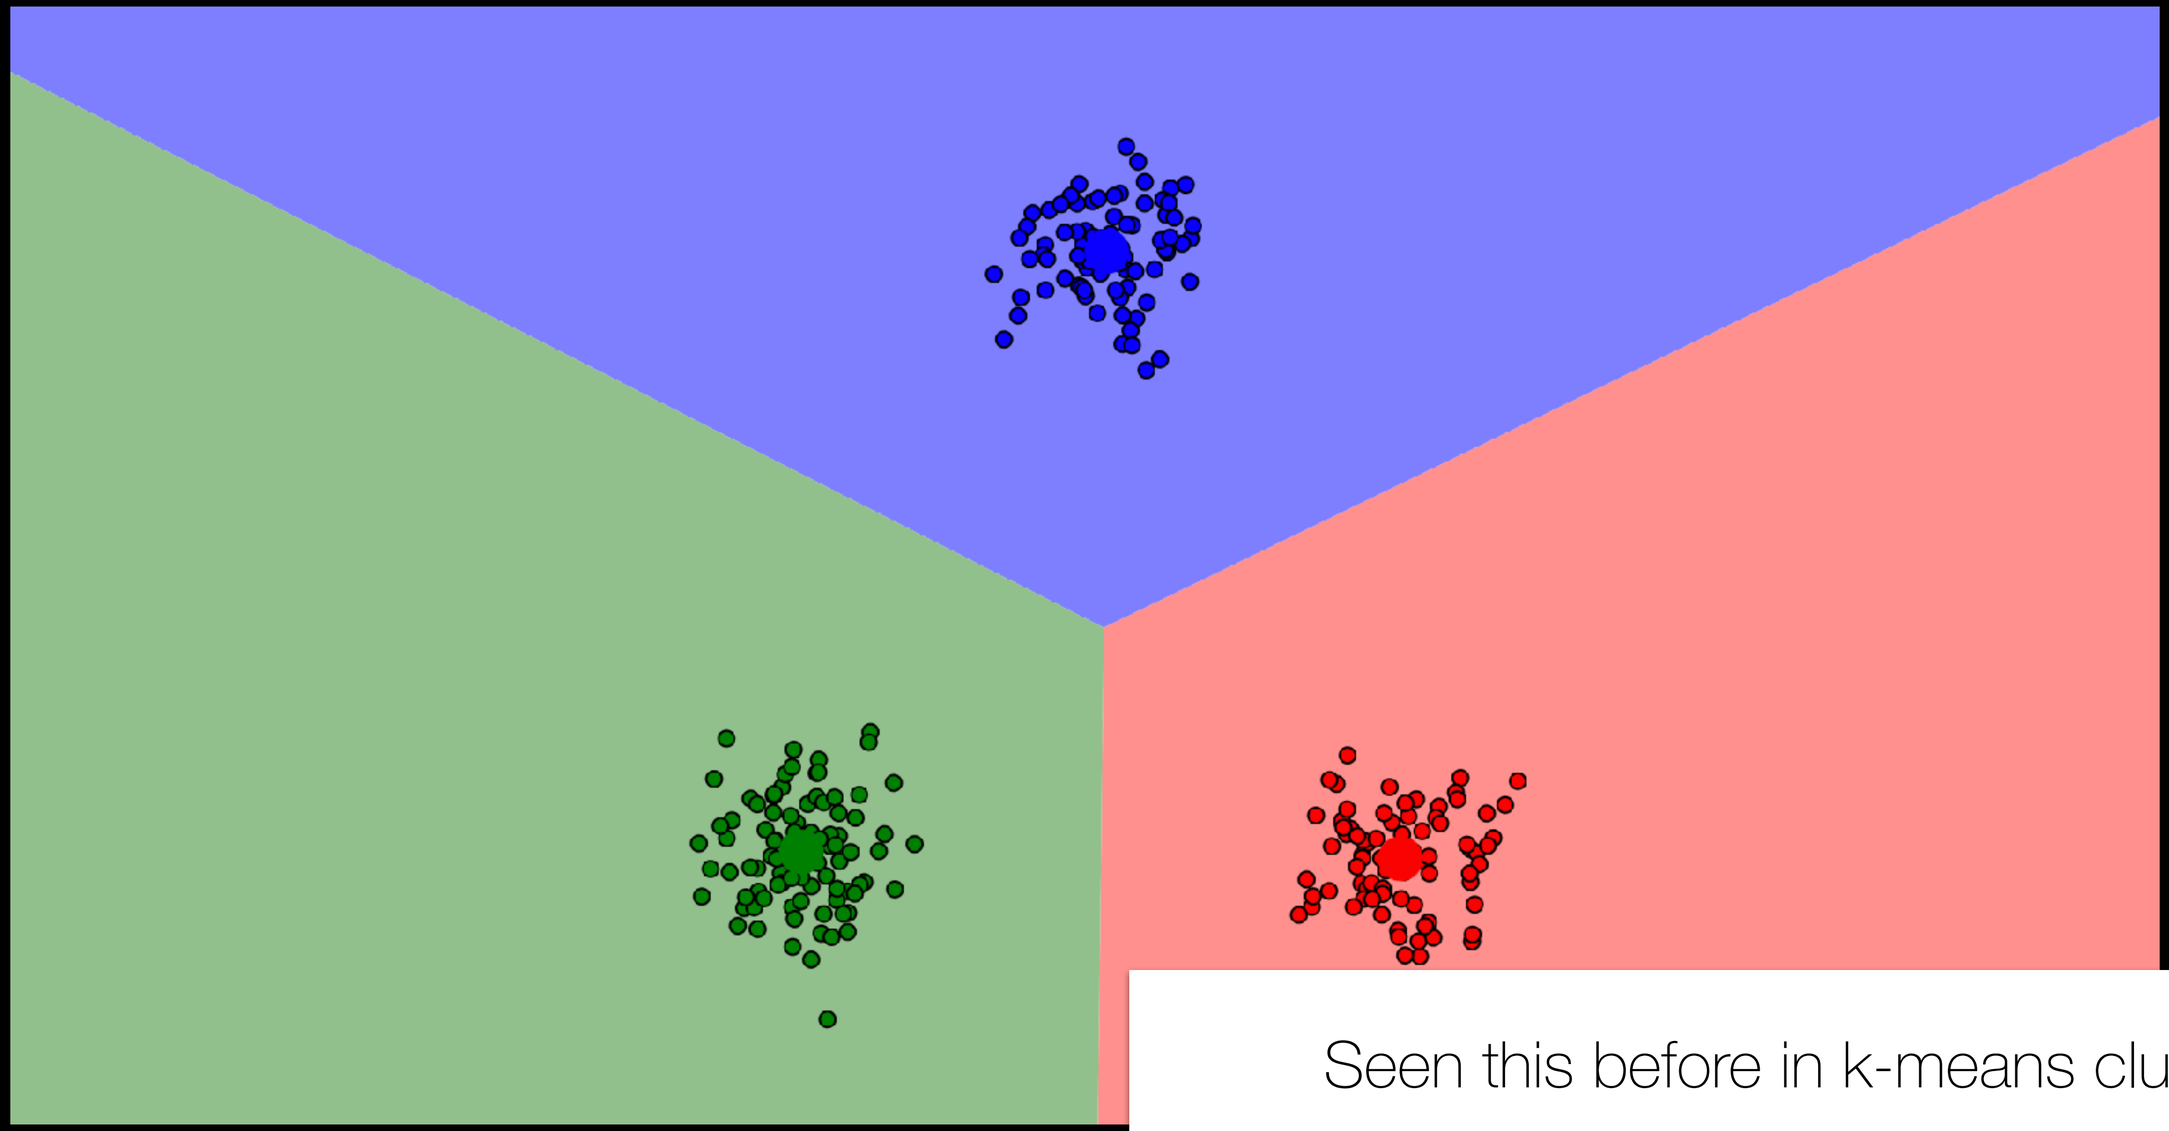

Seen this before in k-means clustering

KNN Classifier

- Model complexity controlled by k (# of nearest neighbors)
- Larger k indicates less complex decision boundary

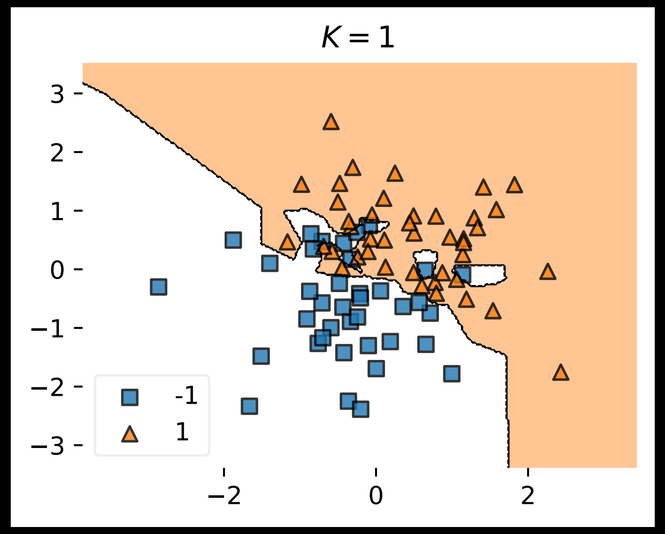

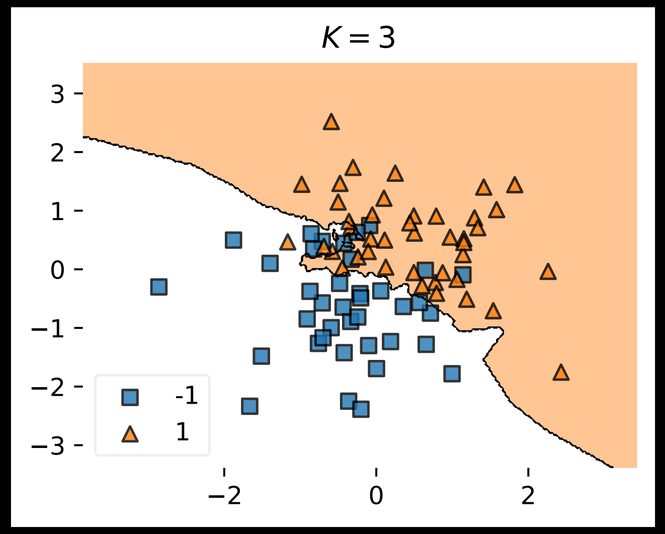

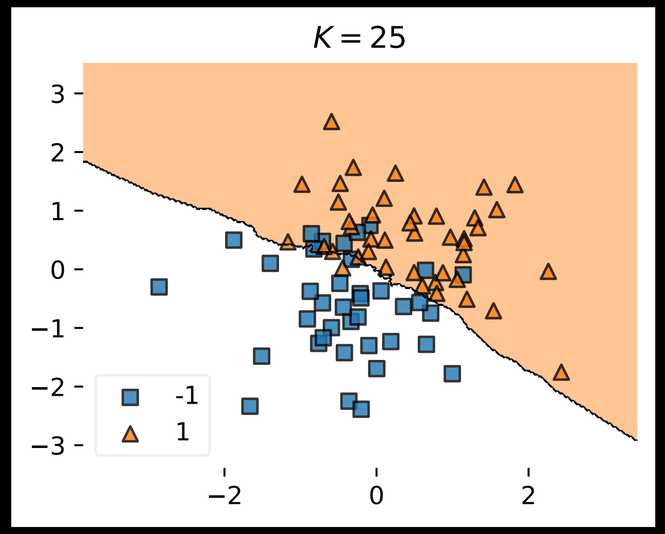

KNN for Regression (example: 1 feature)

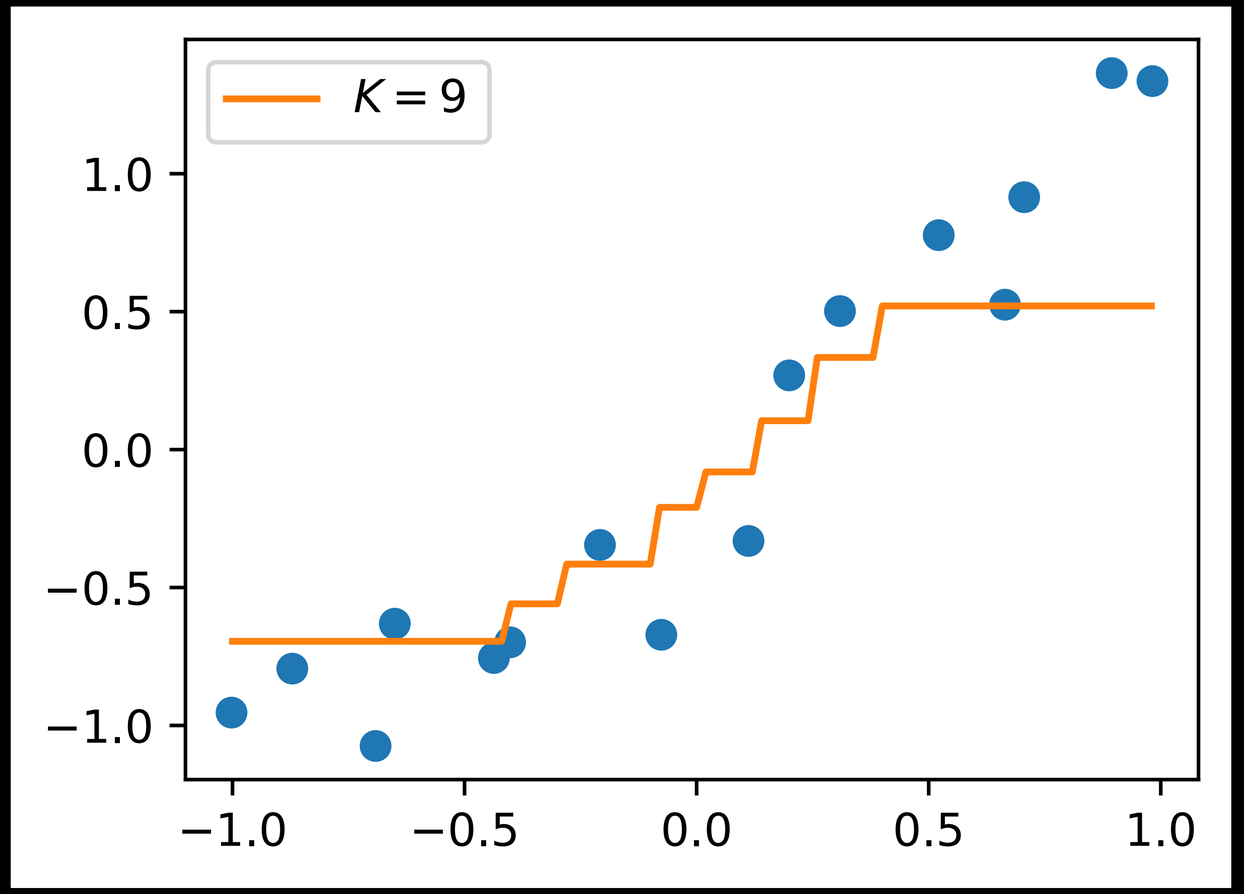

KNN for Regression (example: 1 feature)

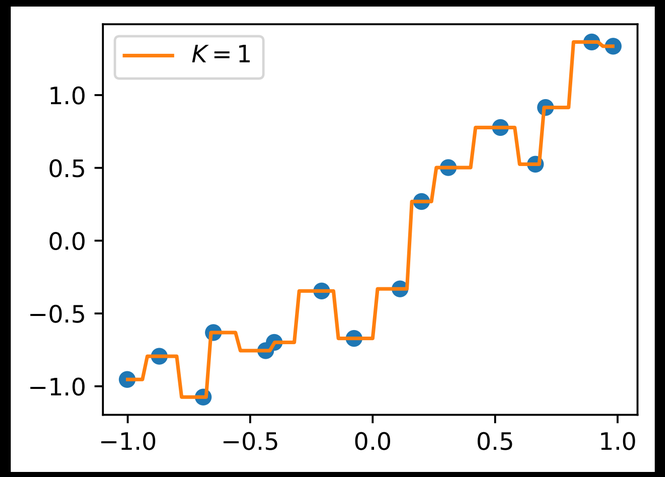

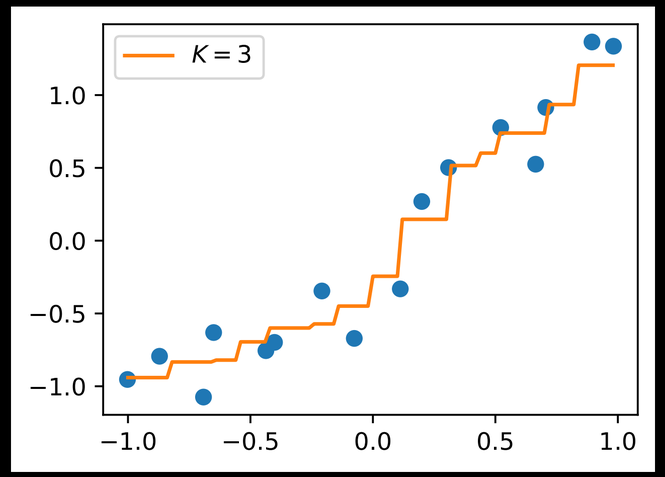

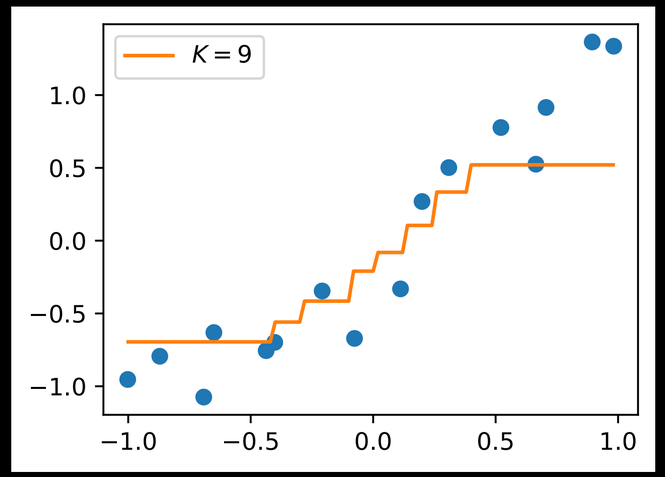

KNN Classifier and Regression

- Decision boundaries become smoother with larger k values
  - In both classification and regression

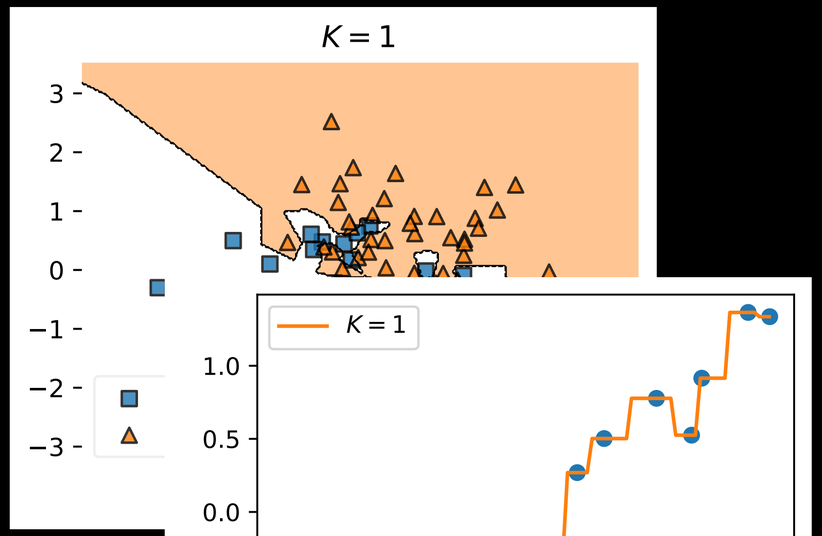

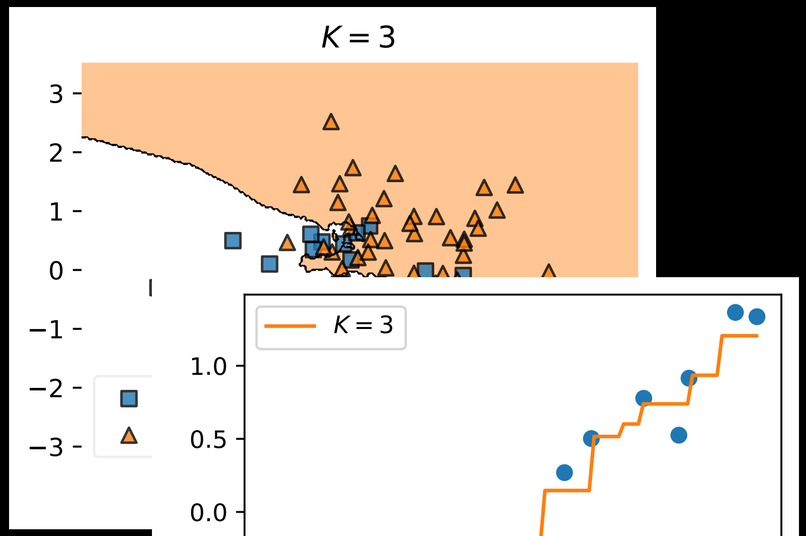

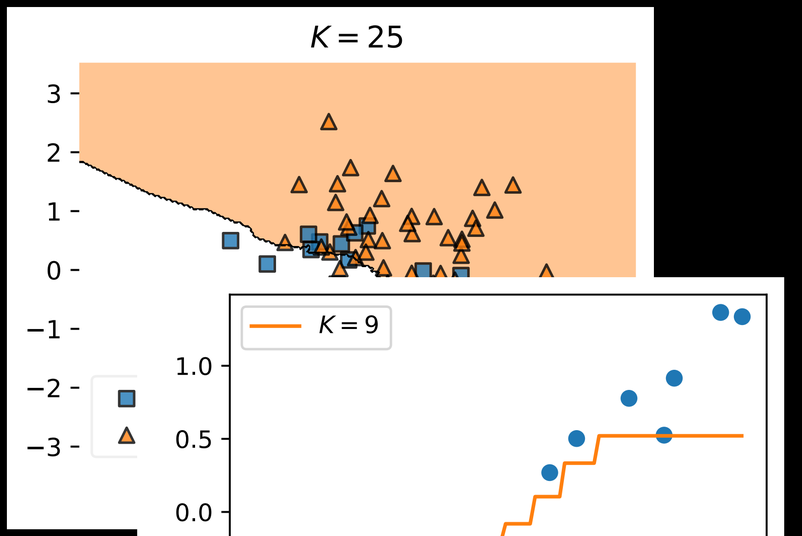

This Module's Learning Objectives

Part 1

- Differentiate between classification and regression in supervised learning
- Describe how voting is used in for k-nearest neighbors classification
- Differentiate binary, multi-class, and multi-label classification
- Define overfitting and describe its impact on generalizability

This Module's Learning Objectives

Part 1

- Differentiate between classification and regression in supervised learning
- Describe how voting is used in for k-nearest neighbors classification
- Differentiate binary, multi-class, and multi-label classification
- Define overfitting and describe its impact on generalizability

Classifying into two labels (yes/no, class 1/2, etc) is super common

But not the only task

Binary Classification

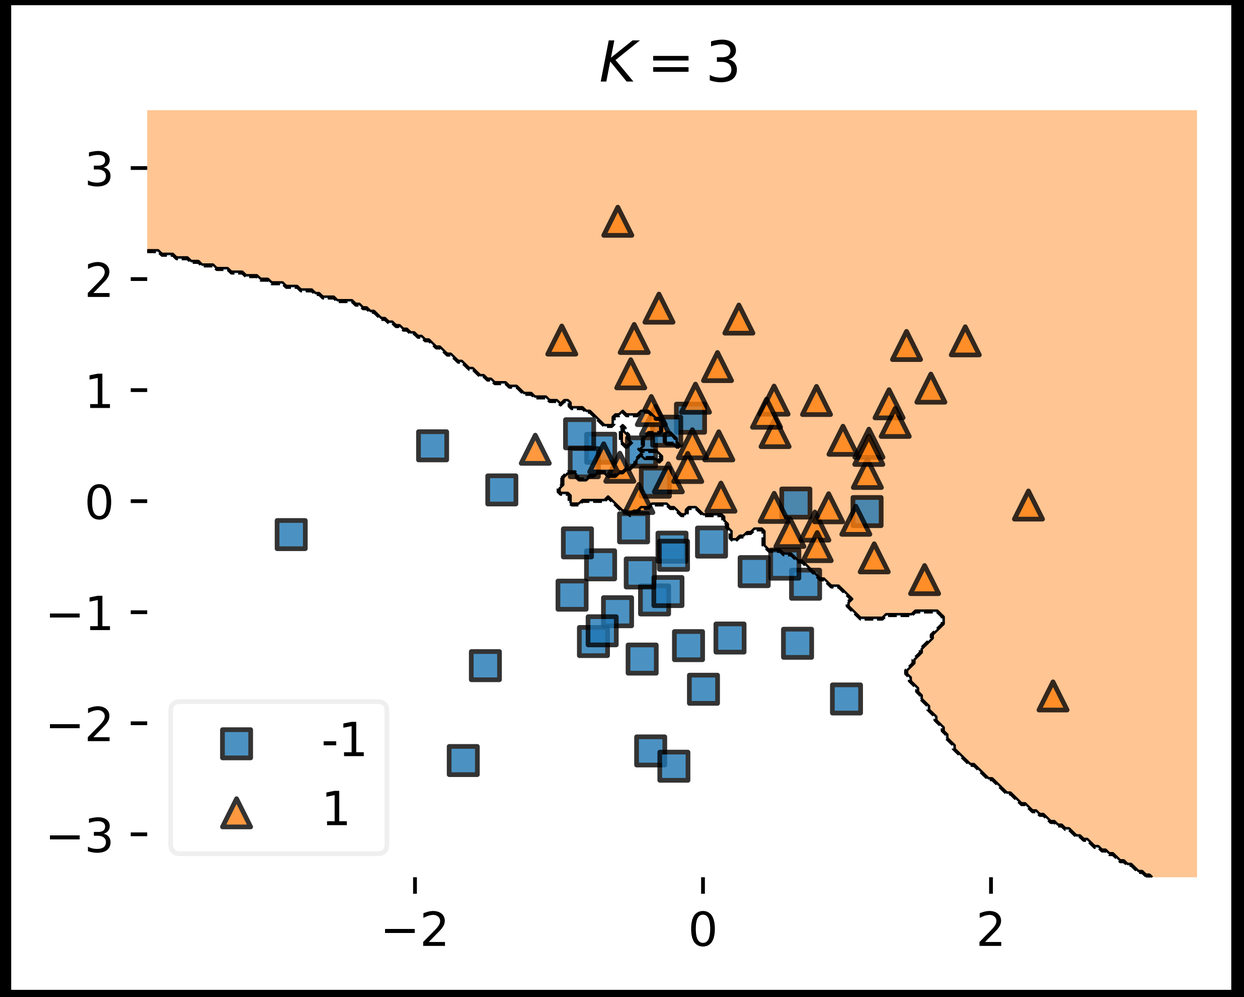

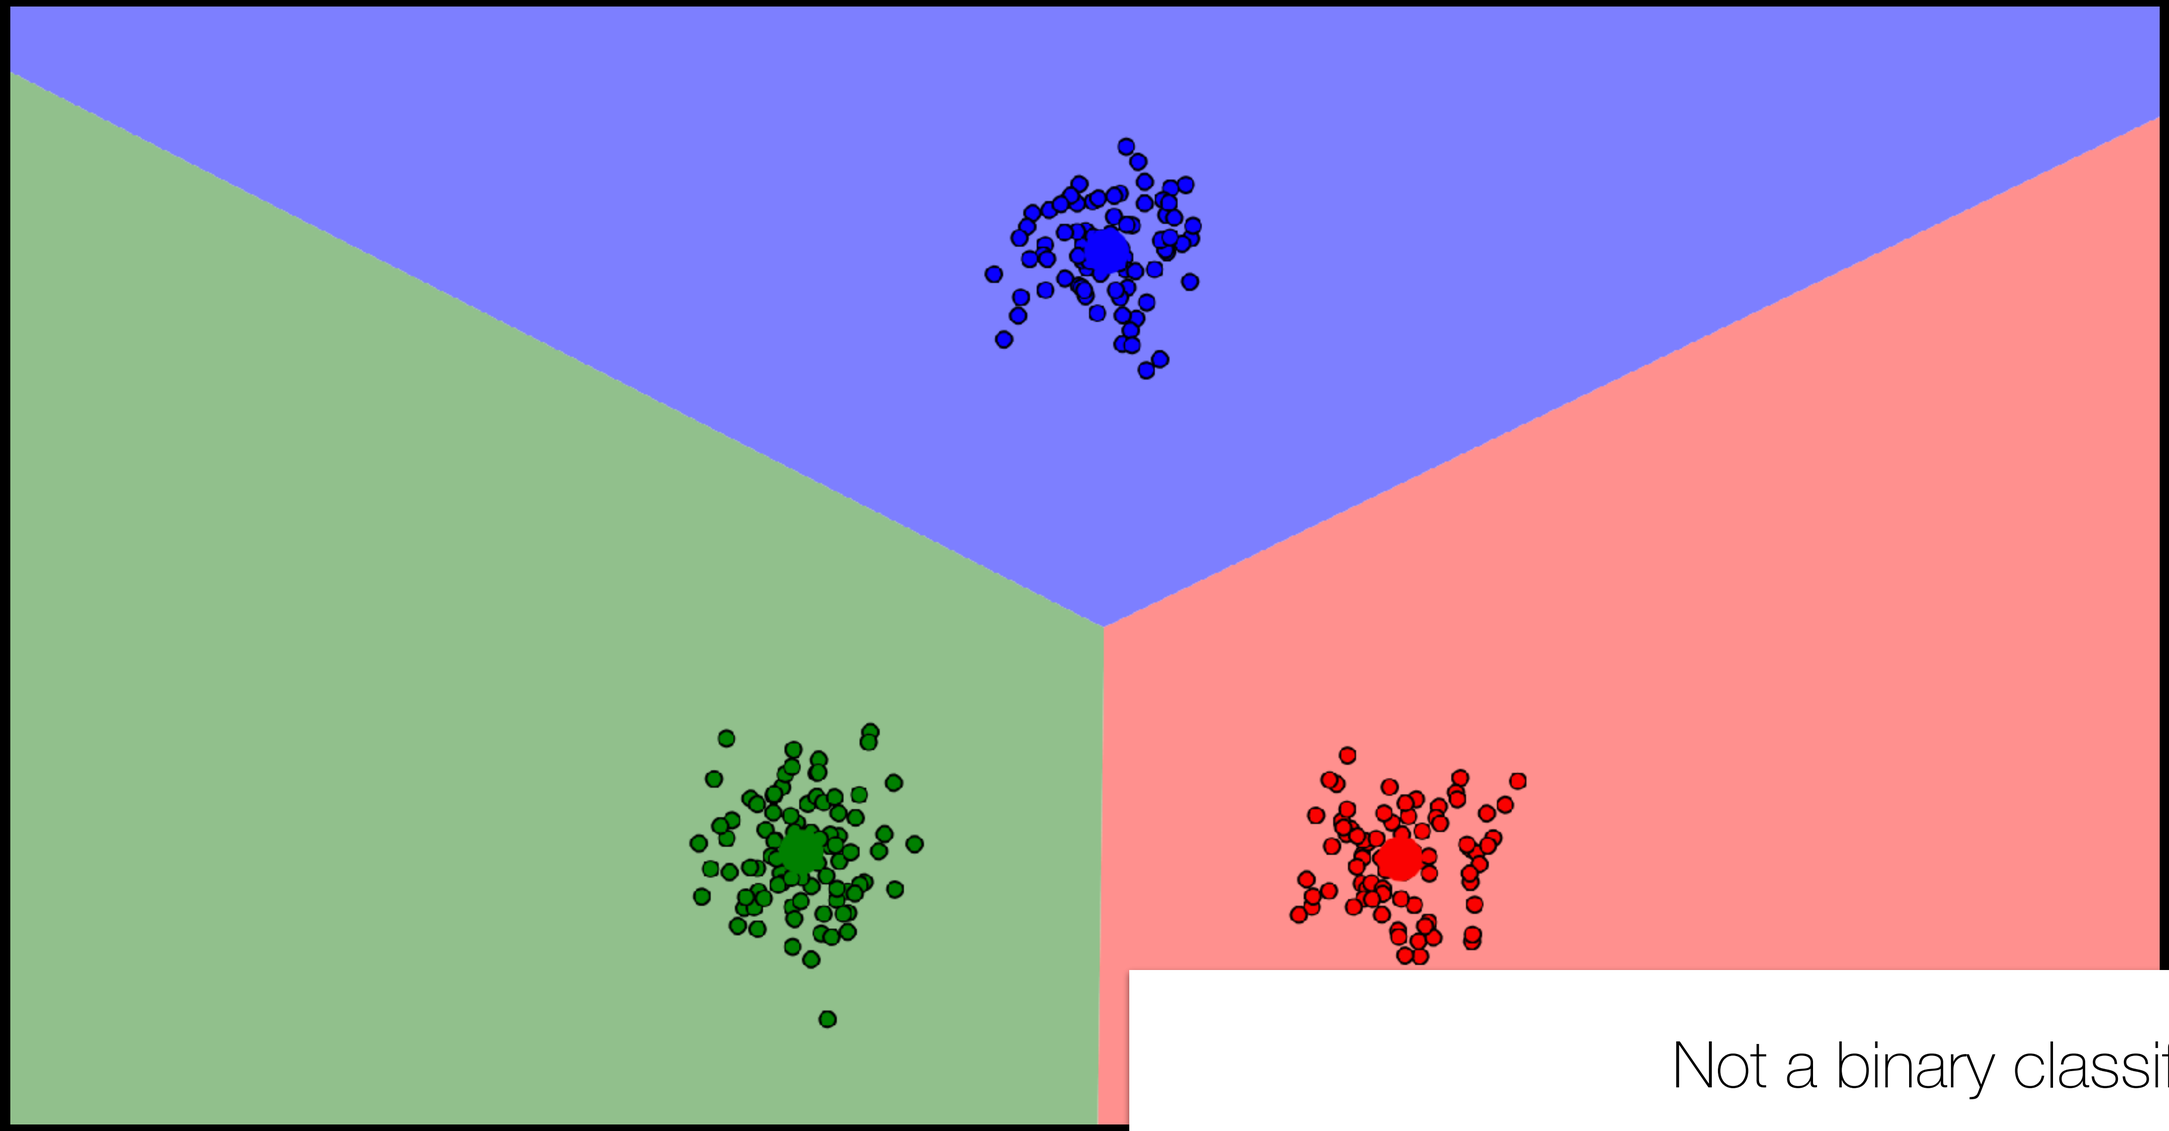

Not a binary classification

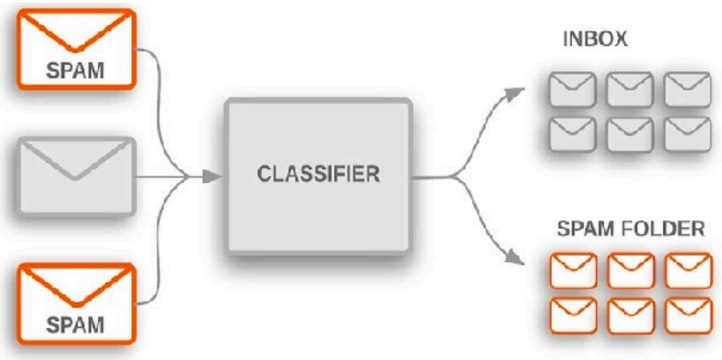

Binary Classification
e.g., Spam vs. Not-Spam

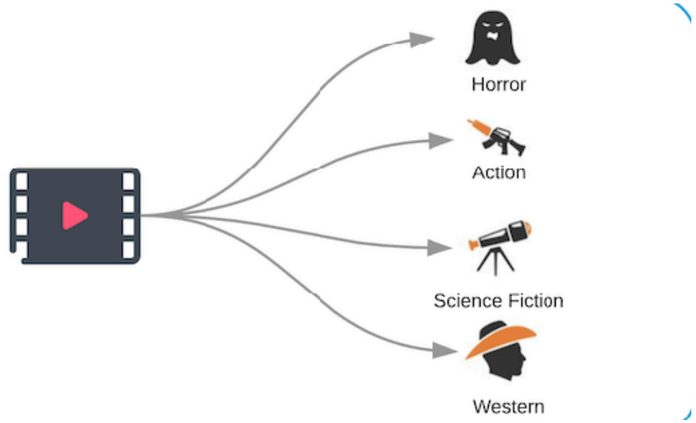

Multi-Class Classification
e.g., Movie Genre

Accuracy in Multi-Class Classification

Can still calculate accuracy here

Accuracy =

$$\frac{1}{|X|} \sum_i f(x_i) == y_i$$

|X| = # of samples

Was the prediction the same as the actual label?

But can’t a movie be in more than one genre?

Can assume only one genre for standard multi-class classification

BUT! We can handle multiple labels

Multi-Label Classification: Samples can have multiple labels associated with them

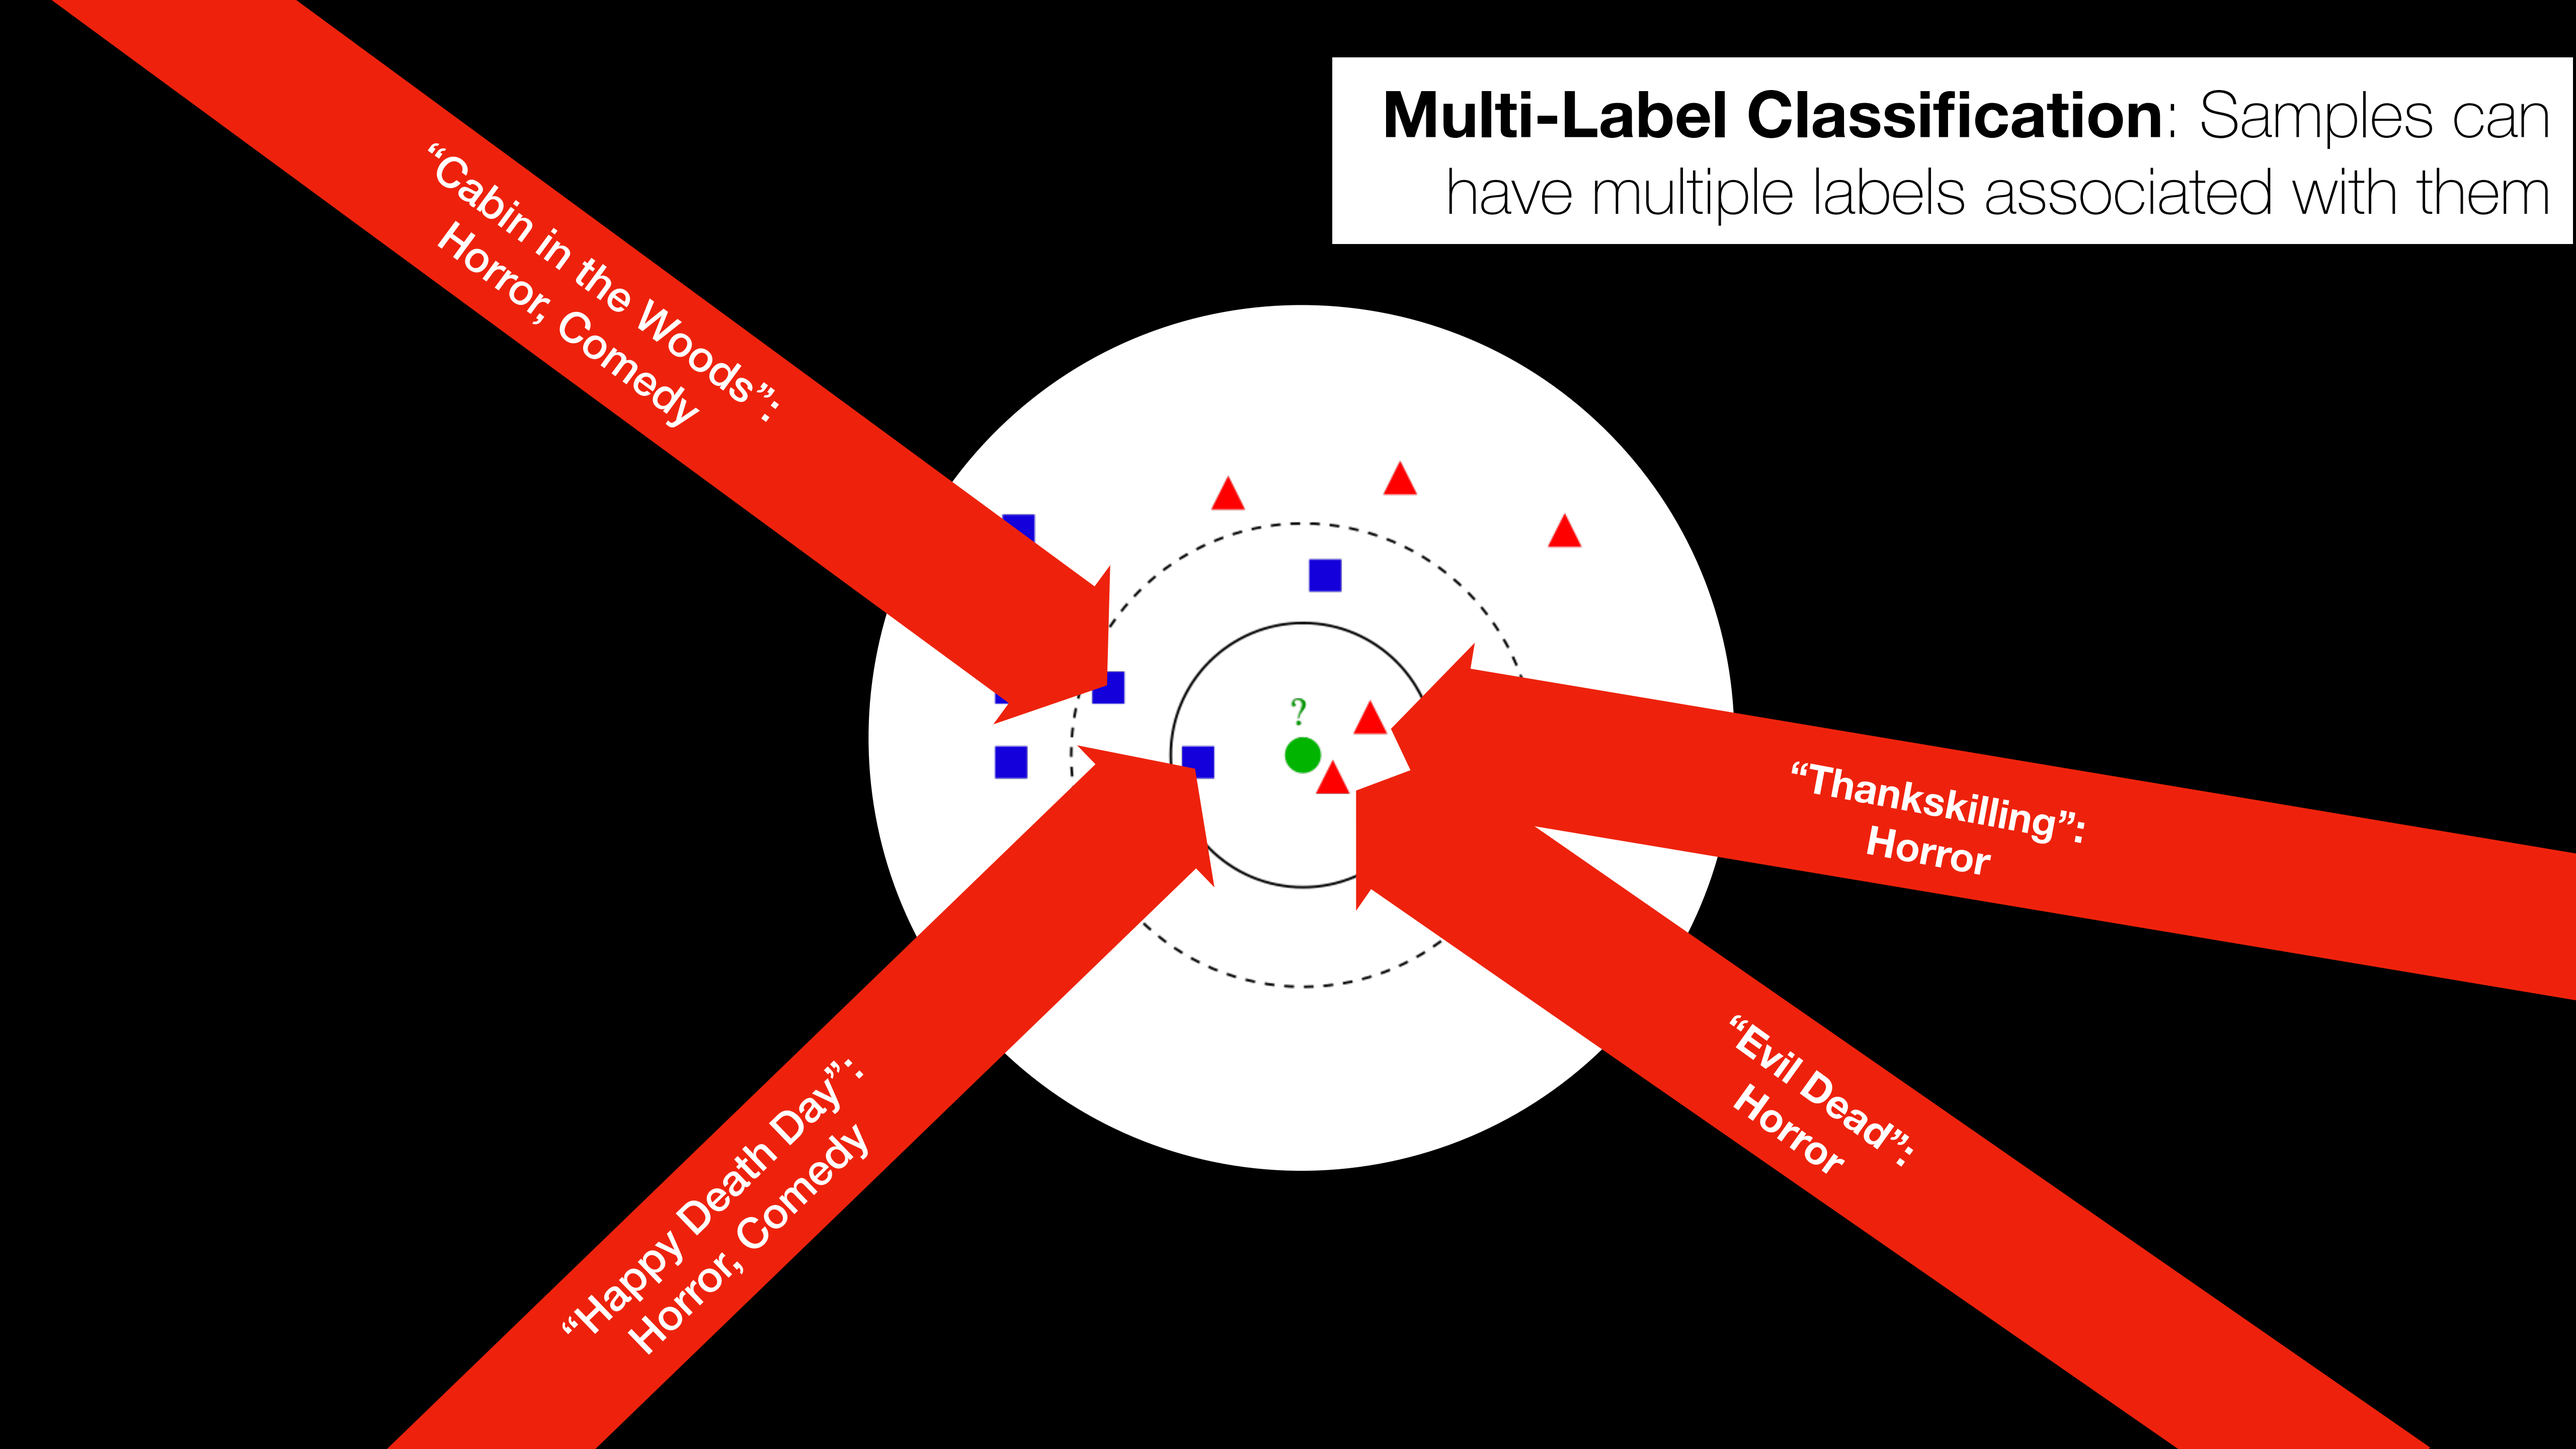

This Module's Learning Objectives

Part 1

- Differentiate between classification and regression in supervised learning
- Describe how voting is used in for k-nearest neighbors classification
- Differentiate binary, multi-class, and multi-label classification
- Define overfitting and describe its impact on generalizability

Today's Exercises

Exercise 1: Predict Genre Based on k Nearest Neighbors

(Extra) Exercise 2: Predict Rating Based on k Nearest Neighbors

Exercise 1: Predict Genre

Exercise 1. Predicting Movie Genre Based on k Nearest Neighbors

1. Download the scaffolding code: Week 11 Exercise-Scaffolding.ipynb
2. For a given target movie, use Jaccard similarity using starring actors to predict the target movie's genre based on the k most similar movies
   A. I.e., calculate Jaccard similarity between two movies using the actors starring in both films, and rank the top-k most similar films
   B. Use the most common genre among these top-k films as the predicted genre for your target movie
3. For each of the following target movies and k=1, determine the most common genre(s) in the most similar movie
   1. The Incredibles (tt0317705)
   2. Interstellar (tt0816692)
   3. The Notebook (tt0332280)
4. Repeat the above using k=3 and k=5. How does the inferred genre change as we increase k?

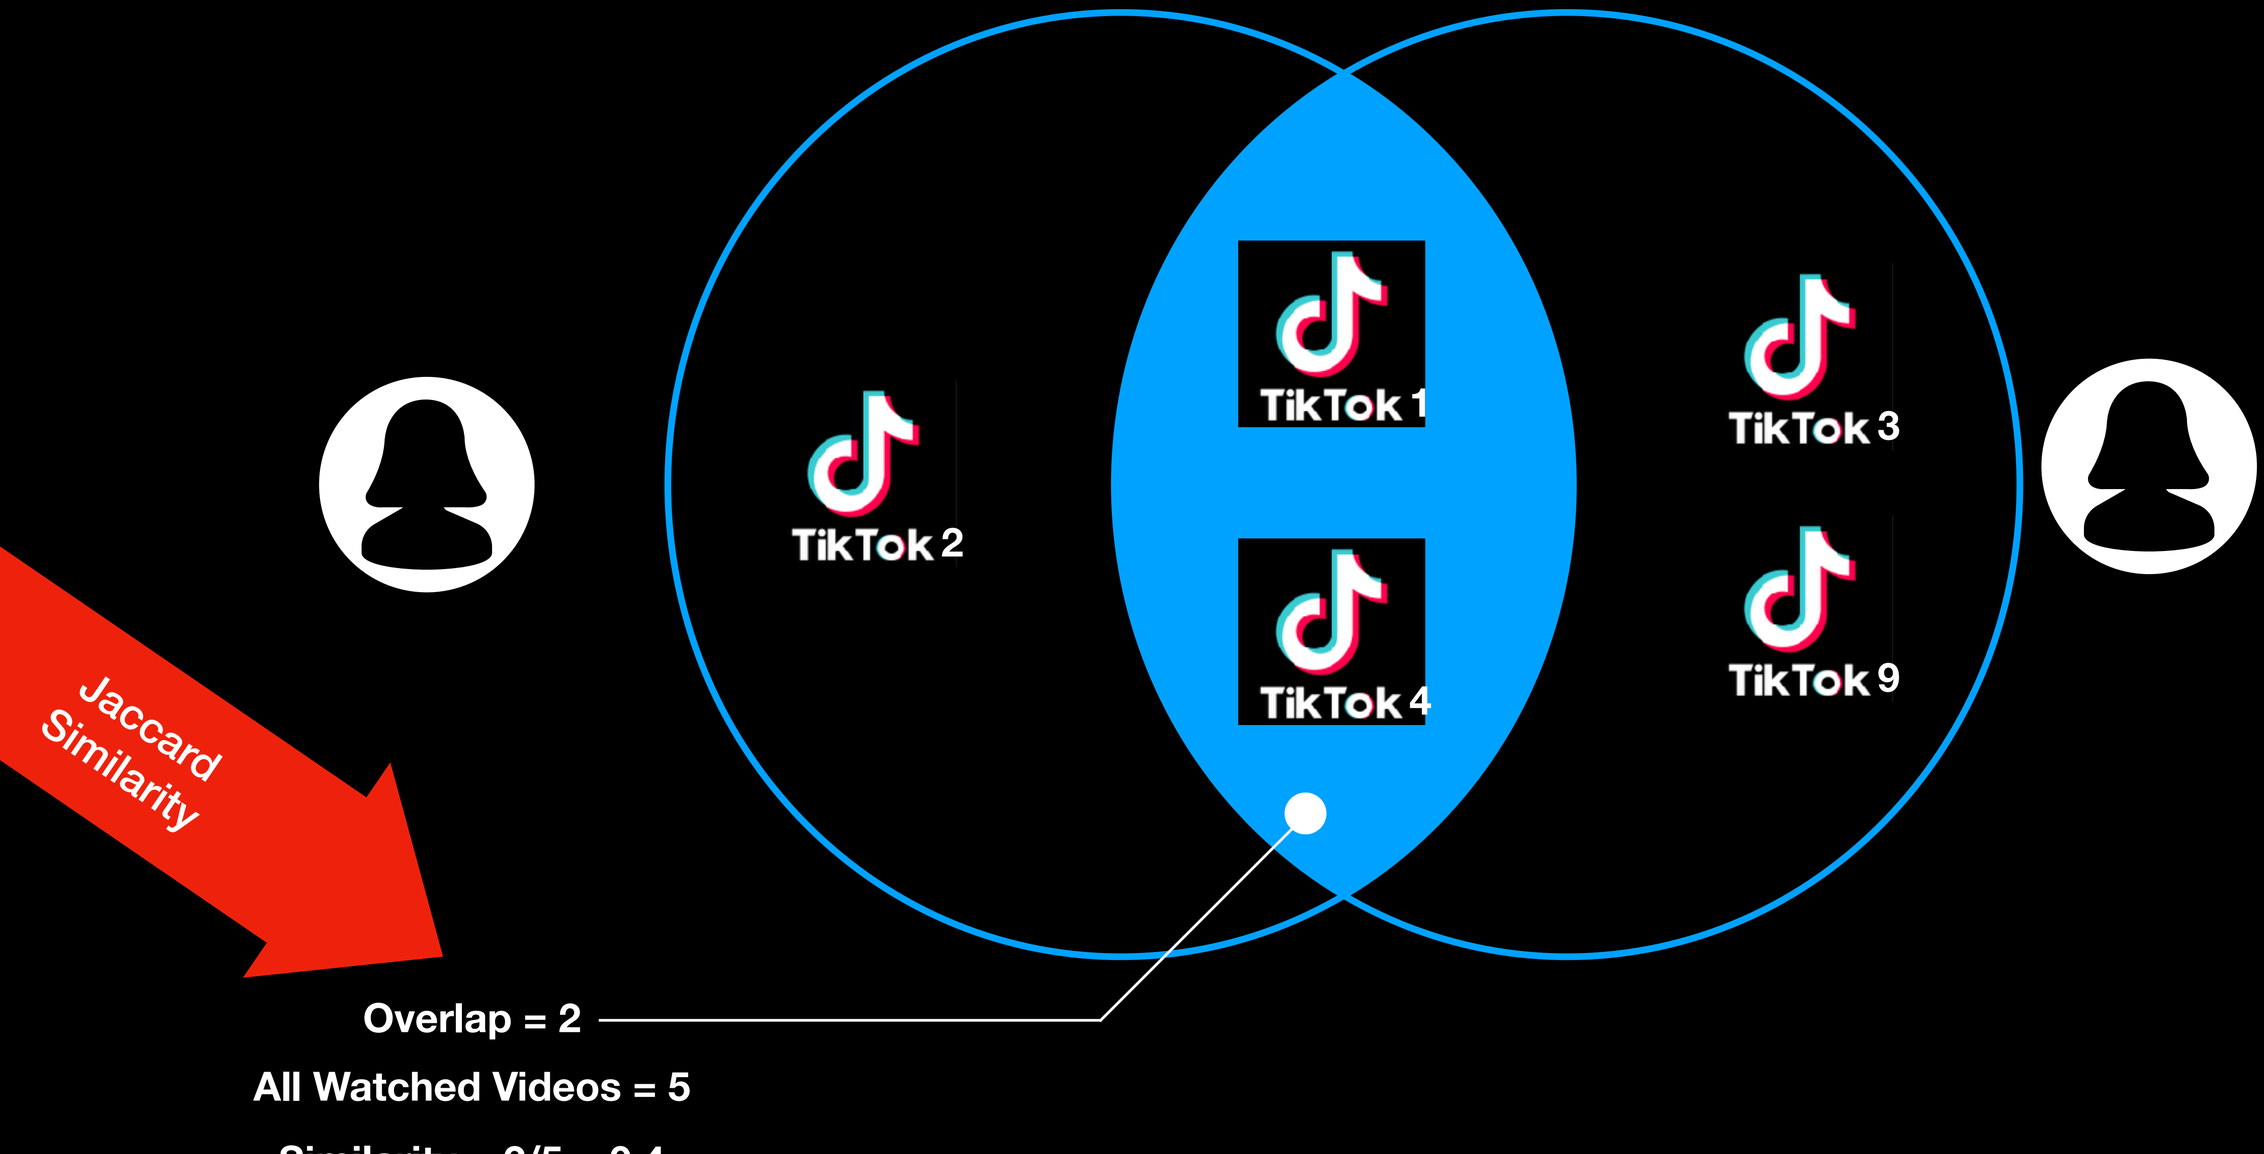
TikTok 1
TikTok 2
TikTok 3
TikTok 4
TikTok 9
Jaccard Similarity
Overlap = 2
All Watched Videos = 5
Similarity = 2/5 = 0.4

(Extra) Exercise 2: Predict Rating

Exercise 2. Predicting Movie Rating Based on k Nearest Neighbors (Extra Practice)

1. Download the scaffolding code: Week 11 Exercise-Scaffolding.ipynb
2. For a given target movie, use Jaccard similarity using starring actors to predict the target movie's rating based on the k most similar movies
   A. I.e., calculate Jaccard similarity between two movies using the actors starring in both films, and rank the top-k most similar films
   B. Use the average rating across these top-k films as the predicted rating for your target movie
3. For each of the following target movies and k=1, determine the most common genre(s) in the most similar movie
   1. The Incredibles (tt0317705)
   2. Interstellar (tt0816692)
   3. The Notebook (tt0332280)
4. Repeat the above using k=3 and k=5. How does the inferred rating change as we increase k?

What are your questions?

Prof. Cody Buntain | @codybuntain | cbuntain@umd.edu
Director, Information Ecosystems Lab