



Supervised Learning, pt 2.2

INST414 - Data Science Techniques

This Module's Learning Objectives

Part 2

Define regularization and its use of improving generalizability

Use linear regression models to predict scores for data elements

Explain how linear regression can be adapted for classification

Extract the class-label probabilities for model outputs

This Module's Learning Objectives

Part 2

Define regularization and its use of improving generalizability

Use linear regression models to predict scores for data elements

Explain how linear regression can be adapted for classification

Extract the class-label probabilities for model outputs

Regularization: Methods for penalizing model complexity

Reducing complexity reduces overfitting

Increases model interpretability

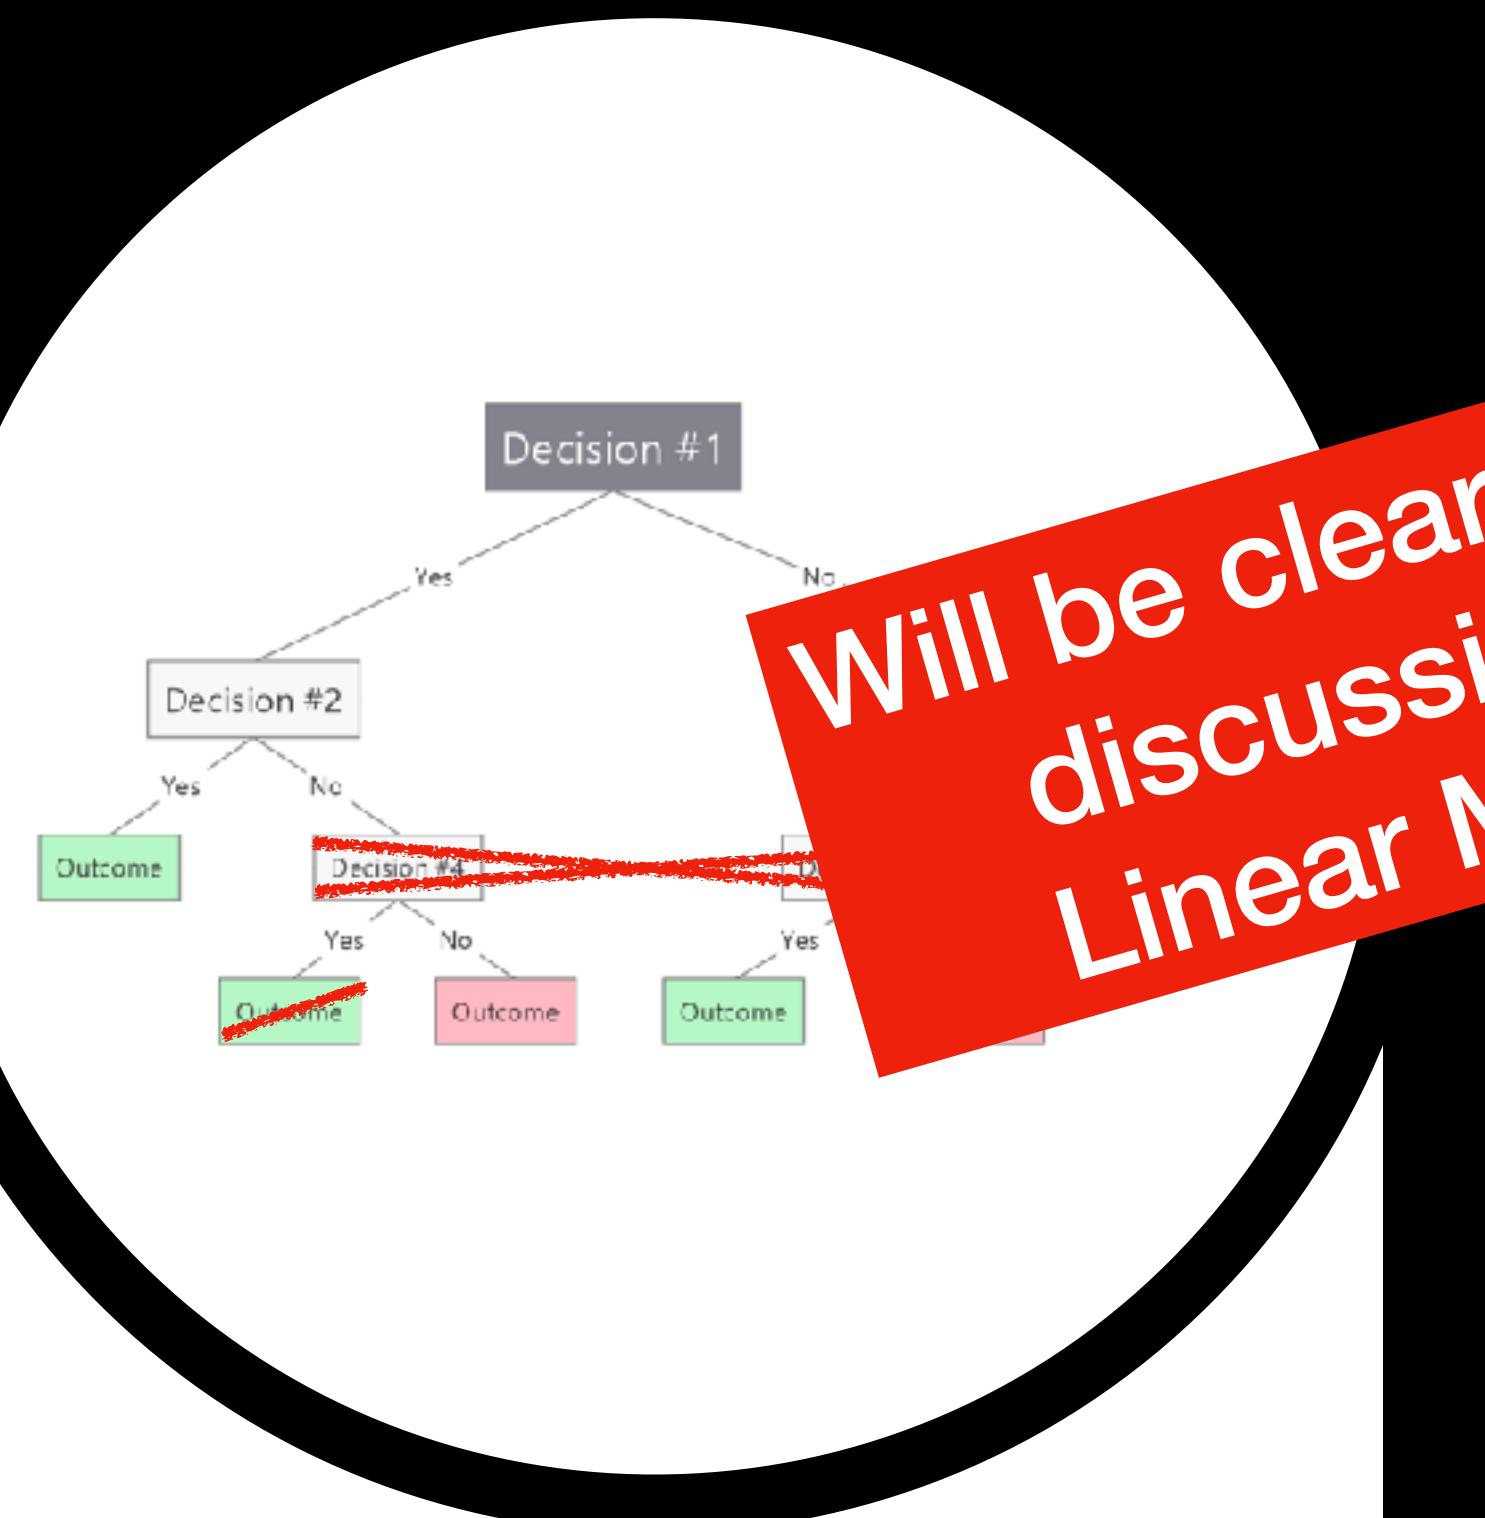
(Often) Decreases computational burden

Regularization: Methods for penalizing model complexity

How might you penalize model complexity?

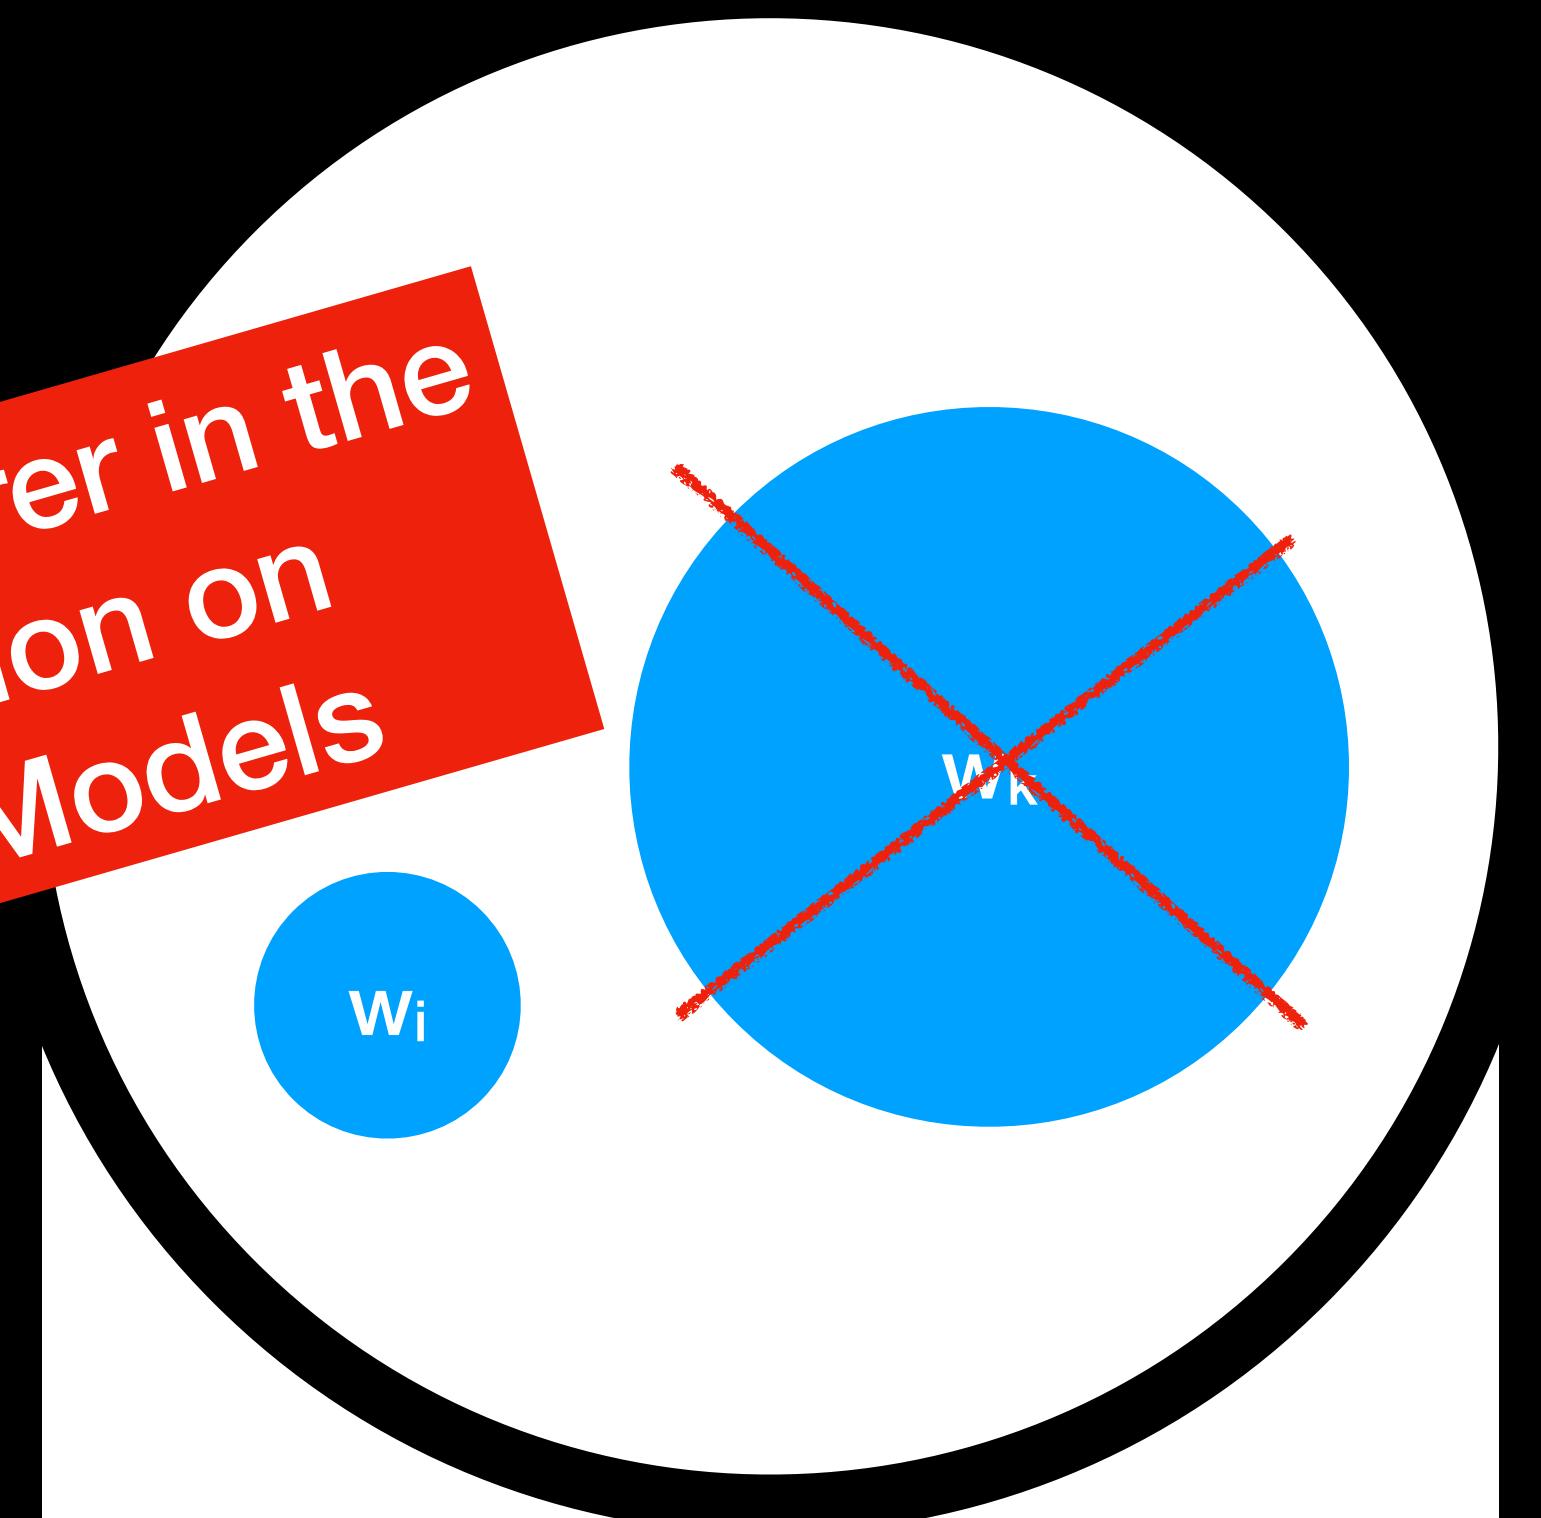
x0	x1	x2	x3	x4	x5
1	1.05	9.2	151	54.4	1.6
2	0.89	10.3	202	57.9	2.2
3	1.43	15.4	111	53	3.4
4	1.02	11.2	163	56	0.3
5	1.49	8.8	112	51.2	1
6	1.32	13.5	111	60	-2.2
7	1.22	12.2	75	67.6	2.2
8	1.1	9.2	245	57	3.3
9	1.34	13	168	60.4	7.2
10	1.12	12.4	197	53	2.7
11	0.75	7.5	173	51.5	6.5
12	1.13	10.9	178	62	3.7
13	1.15	12.7	199	53.7	6.4
14	1.09	12	96	49.8	1.4
15	0.96	7.6	164	62.2	-0.1
16	1.16	9.9	252	56	9.2
17	0.76	6.4	136	61.9	9
18	1.05	12.6	150	56.7	2.7
19	1.16	11.7	104	54	-2.1
20	1.2	11.8	148	59.9	3.7
21	1.04	8.6	204	61	
22	1.07	9.3	174	54.3	

You use fewer dimensions/features



Will be clearer in the discussion on Linear Models

Model enforces sparsity (mostly 0s)



Model forces feature weights to be small

This Module's Learning Objectives

Part 2

Define regularization and its use of improving generalizability

Use linear regression models to predict scores for data elements

Explain how linear regression can be adapted for classification

Extract the class-label probabilities for model outputs

Recall from last time...

X y
 \mathcal{M}

mass	width	height	color_score	fruit_label
192	8.4	7.3	0.55	1
180	8.0	6.8	0.59	1
176	7.4	7.2	0.60	1
86	6.2	4.7	0.80	2
84	6.0	4.6	0.79	2
80	5.8	4.3	0.77	2
80	5.9	4.3	0.81	2
76	5.8	4.0	0.81	2
178	7.1	7.8	0.92	1
172	7.4	7.0	0.89	1
166	6.9	7.3	0.93	1
172	7.1	7.6	0.92	1
154	7.0	7.1	0.88	1
164	7.3	7.7	0.70	1
152	7.6	7.3	0.69	1
156	7.7	7.1	0.69	1
156	7.6	7.5	0.67	1
168	7.5	7.6	0.73	1
162	7.5	7.1	0.83	1
162	7.4	7.2	0.85	1

$x_i \Rightarrow$ ith row in \mathbf{X}

$x_{ij} \Rightarrow$ jth feature of the ith row in \mathbf{X}

\mathbf{X} is an $n \times d$ matrix, where d is the number of features

Linear Regression in simplest form

Linear regression
is the process
of finding these
weights w_j such
that...

Function of weights w times a row in X plus some constant value

We minimize
this error

$$= wx_i + b \mid x_i \in X, y_i \in y$$

Actual Value: y_i , Predicted Value: \hat{y}_i

$$\min(y_i - \hat{y}_i)$$

Linear Regression: Assumes relationship between X and y is linear

May (likely) not be true

Linear Regression is “easy” to interpret

Feature weights w_j
tell you how much
that feature
contributes to the
outcome

$$y_i = \sum_{j \in d} w_j x_{i,j} + b$$

Please [cite us](#) if you use the software.

[sklearn.linear_model.LinearRegression](#)
Examples using [sklearn.linear_model.LinearRegression](#)

sklearn.linear_model.LinearRegression

```
class sklearn.linear_model.LinearRegression(*, fit_intercept=True, normalize='deprecated', copy_X=True, n_jobs=None, positive=False)
```

[\[source\]](#)

Ordinary least squares Linear Regression.

LinearRegression fits a linear model with coefficients $w = (w_1, \dots, w_p)$ to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation.

Parameters:

fit_intercept : bool, default=True

Whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (i.e. data is expected to be centered).

normalize : bool, default=False

This parameter is ignored when `fit_intercept` is set to False. If True, the regressors X will be normalized before regression by subtracting the mean and dividing by the L2-norm. If you wish to standardize, please use [StandardScaler](#) before calling `fit` on an estimator with `normalize=False`.

Deprecated since version 1.0: `normalize` was deprecated in version 1.0 and will be removed in 1.2.

copy_X : bool, default=True

If True, X will be copied; else, it may be overwritten.

n_jobs : int, default=None

The number of jobs to use for the computation. This will only provide speedup in case of sufficiently large problems, that is if firstly `n_targets > 1` and secondly X is sparse or if `positive` is set to `True`. `None` means 1 unless in a [joblib.parallel_backend](#) context. `-1` means using all processors. See [Glossary](#) for more details.

positive : bool, default=False

When set to `True`, forces the coefficients to be positive. This option is only supported for dense arrays.

New in version 0.24.

Attributes:

coef_ : array of shape (n_features,) or (n_targets, n_features)

Estimated coefficients for the linear regression problem. If multiple targets are passed during the fit (y 2D), this is a 2D array of shape $(n_targets, n_features)$, while if only one target is passed, this is a 1D array of length $n_features$.

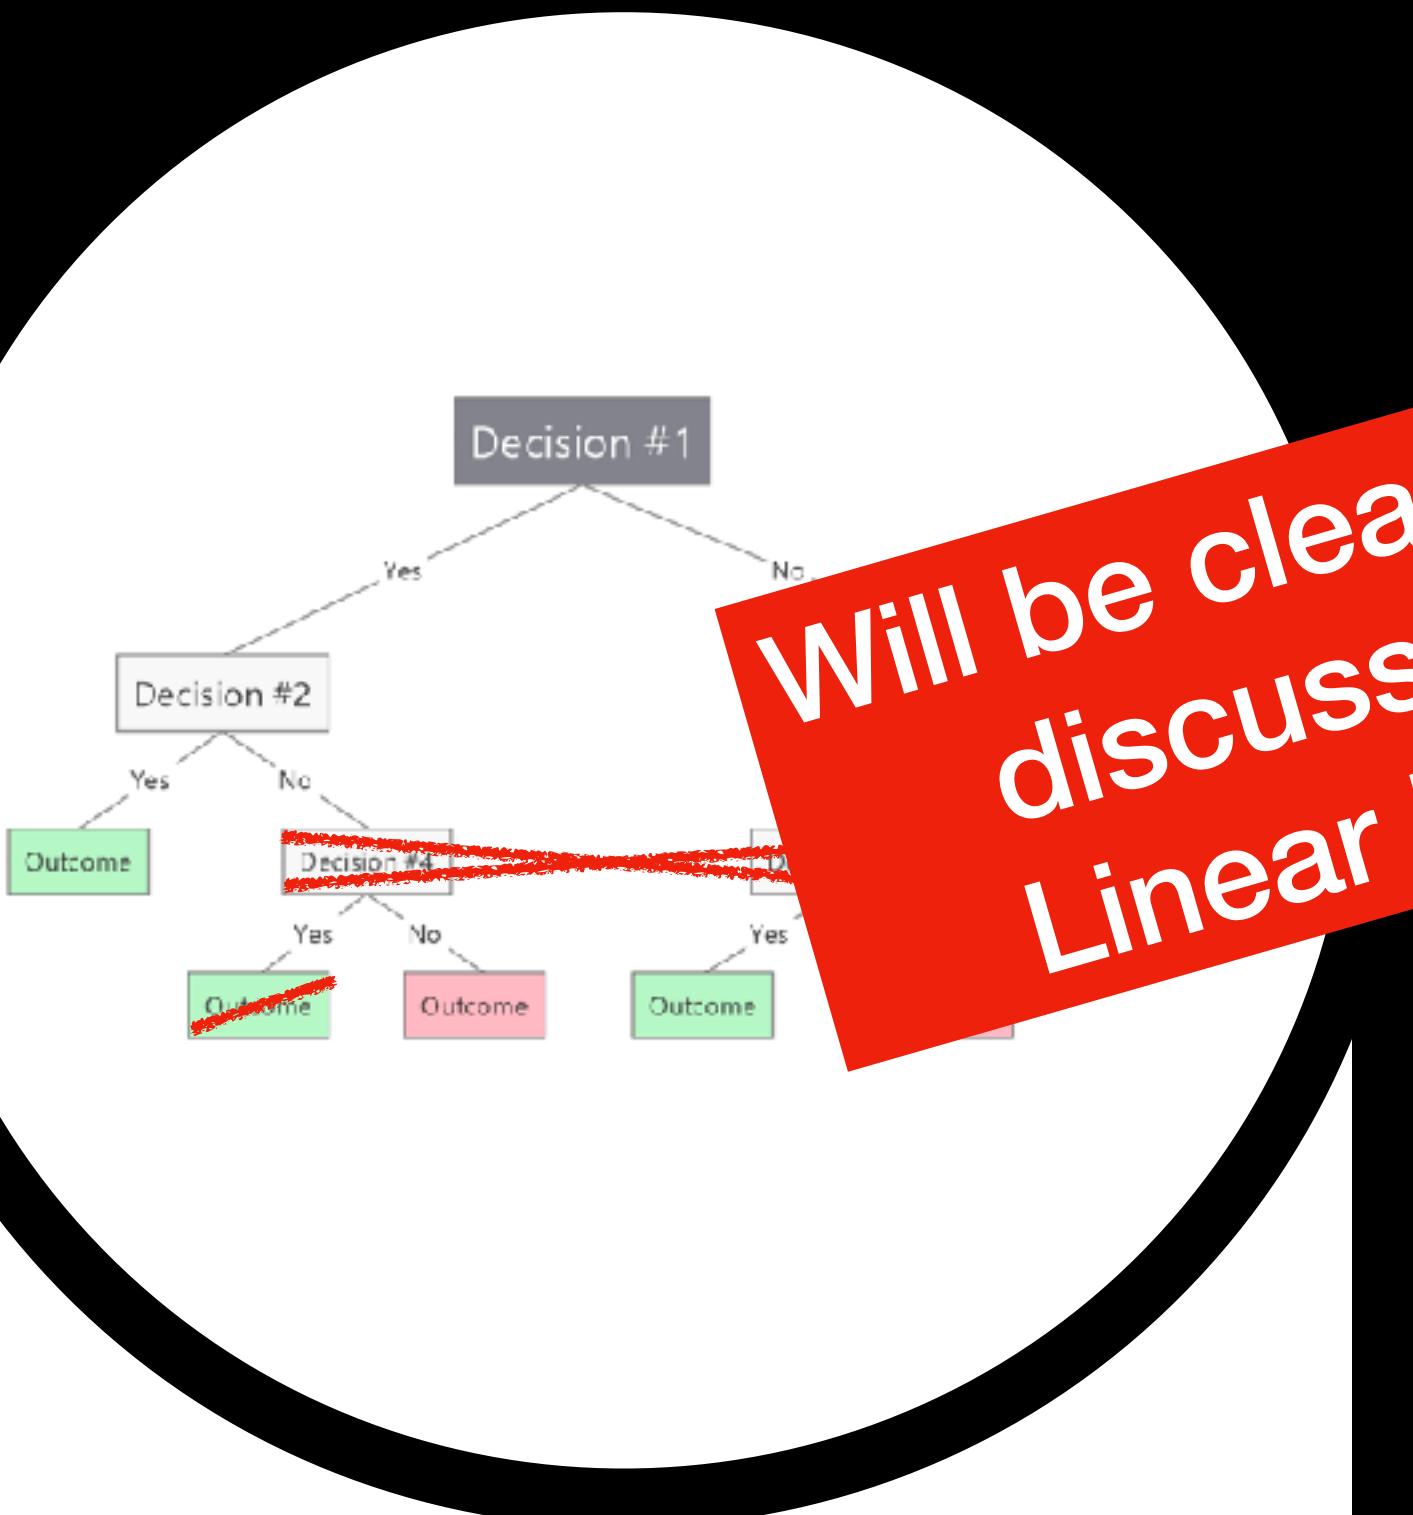
Feature weights w_j tell you how much that feature contributes to the outcome

Regularization: Methods for penalizing model complexity

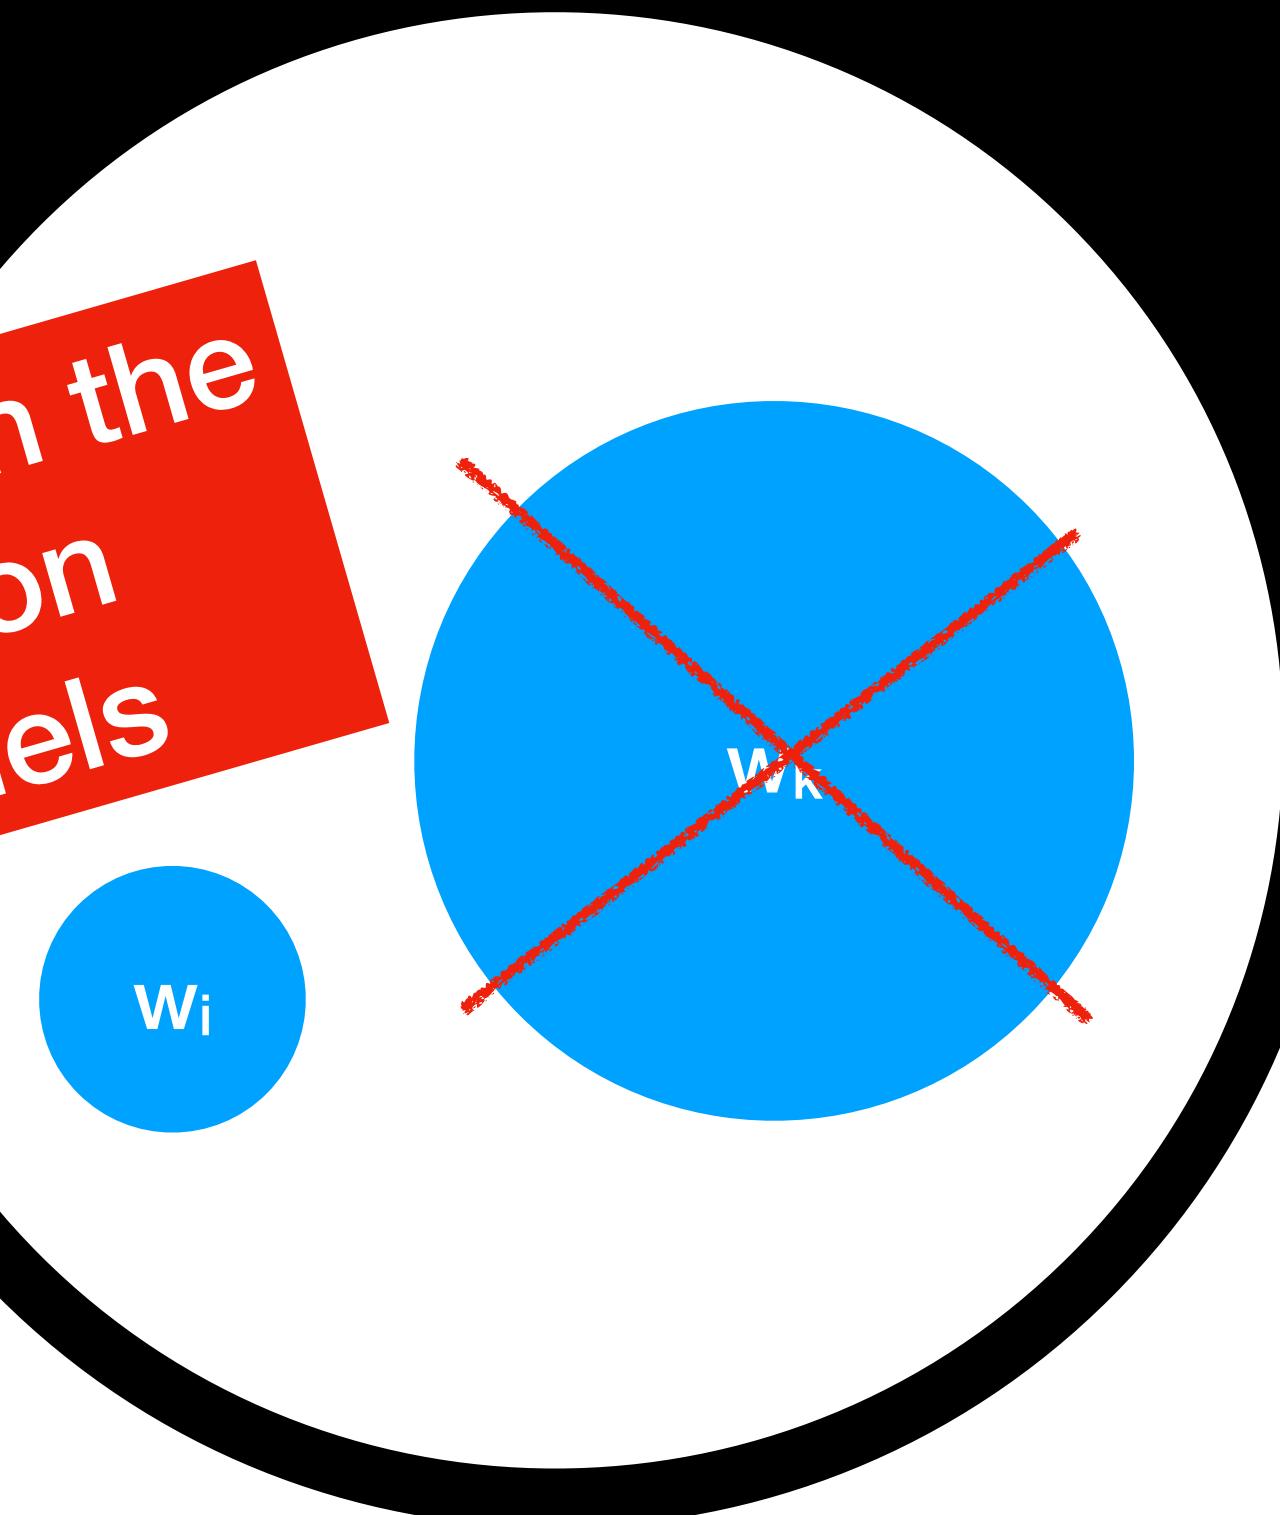
How might you penalize model complexity?

x0	x1	x2	x3	x4	x5
1	1.06	9.2	151	54.4	1.6
2	0.89	10.3	202	57.9	2.2
3	1.43	15.4	113	53	3.4
4	1.02	11.2	163	56	0.3
5	1.49	8.8	192	51.2	1
6	1.32	13.5	111	60	-2.2
7	1.22	12.2	175	67.6	2.2
8	1.1	9.2	245	57	3.3
9	1.34	13	168	60.4	7.2
10	1.12	12.4	197	53	2.7
11	0.75	7.5	173	51.5	6.5
12	1.13	10.9	178	62	3.7
13	1.15	12.7	199	53.7	6.4
14	1.09	12	96	49.8	1.4
15	0.96	7.6	164	62.2	-0.1
16	1.16	9.9	252	56	9.2
17	0.76	6.4	136	61.9	9
18	1.05	12.6	150	56.7	2.7
19	1.16	11.7	104	54	-2.1
20	1.2	11.8	148	59.9	3.7
21	1.04	8.6	204	61	
22	1.07	9.3	174	54.3	

You use fewer dimensions/features



**Model enforces
sparsity (mostly 0s)**



Model forces feature weights to be small

This Module's Learning Objectives

Part 2

Define regularization and its use of improving generalizability

Use linear regression models to predict scores for data elements

Explain how linear regression can be adapted for classification

Extract the class-label probabilities for model outputs

“Transform” Linear Regression for Classification

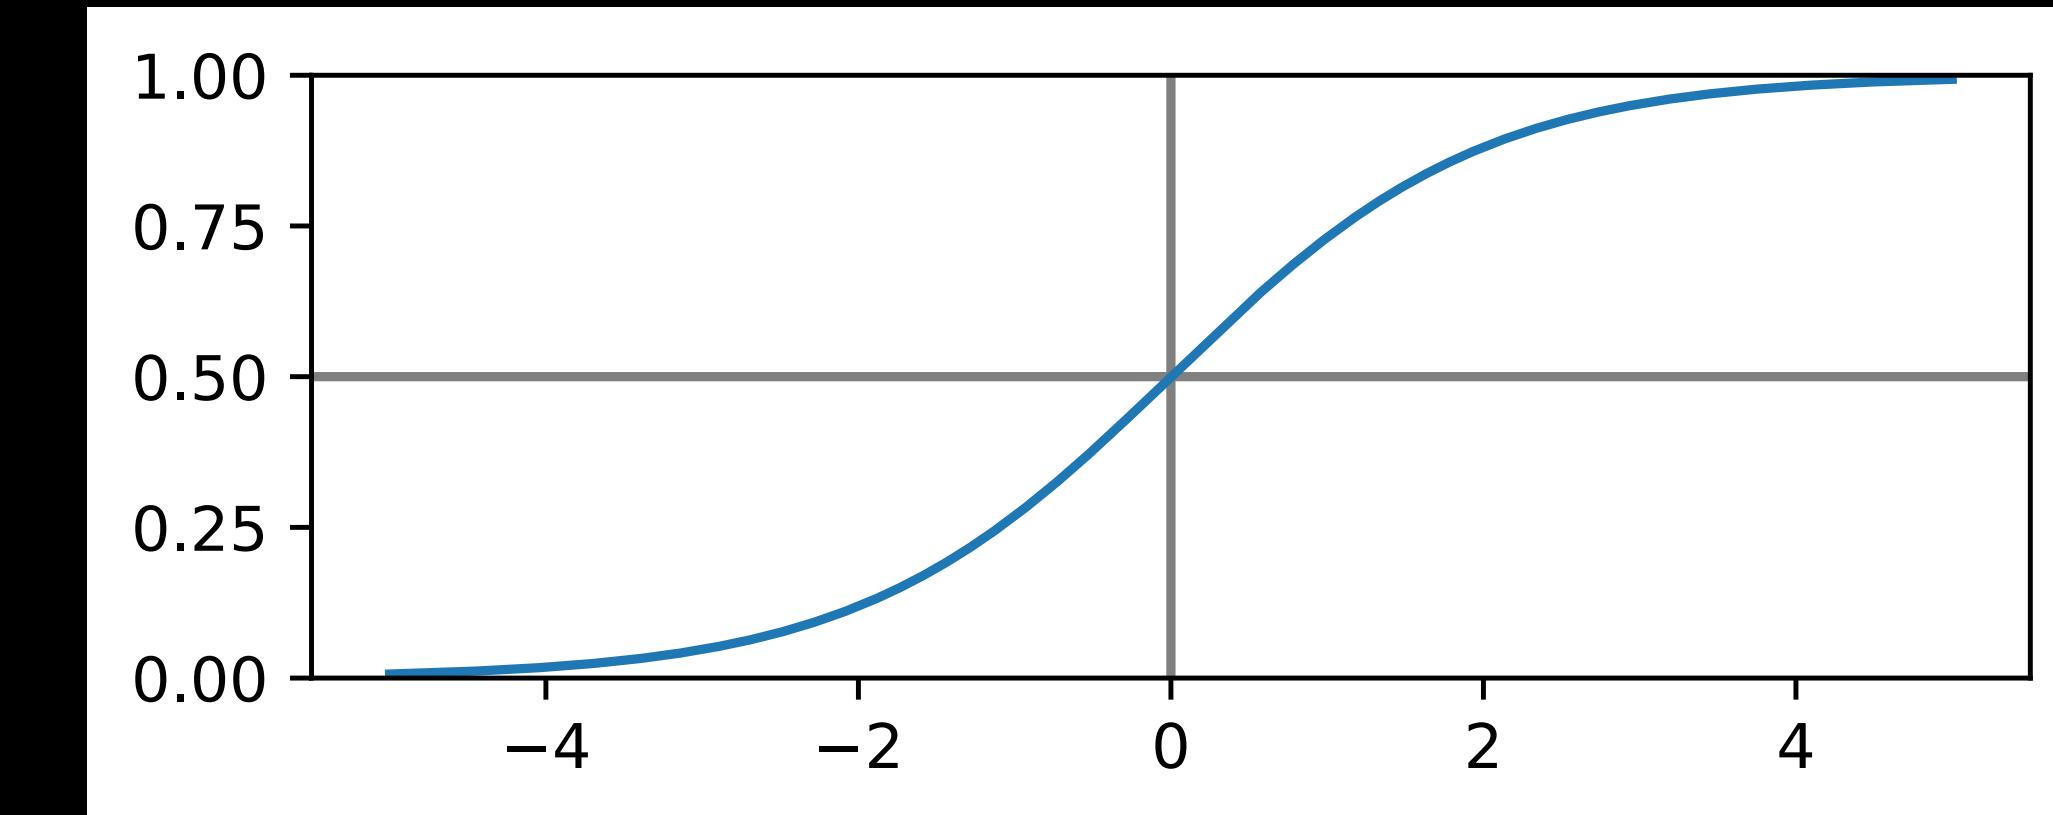
- Regression: $X \rightarrow y, y \in (-\infty, +\infty)$
- Binary classification: $X \rightarrow y, y \in \{0,1\}$
- Can we construct a mapping of $(-\infty, +\infty) \rightarrow \{0,1\}$?

“Transform” Linear Regression for Classification

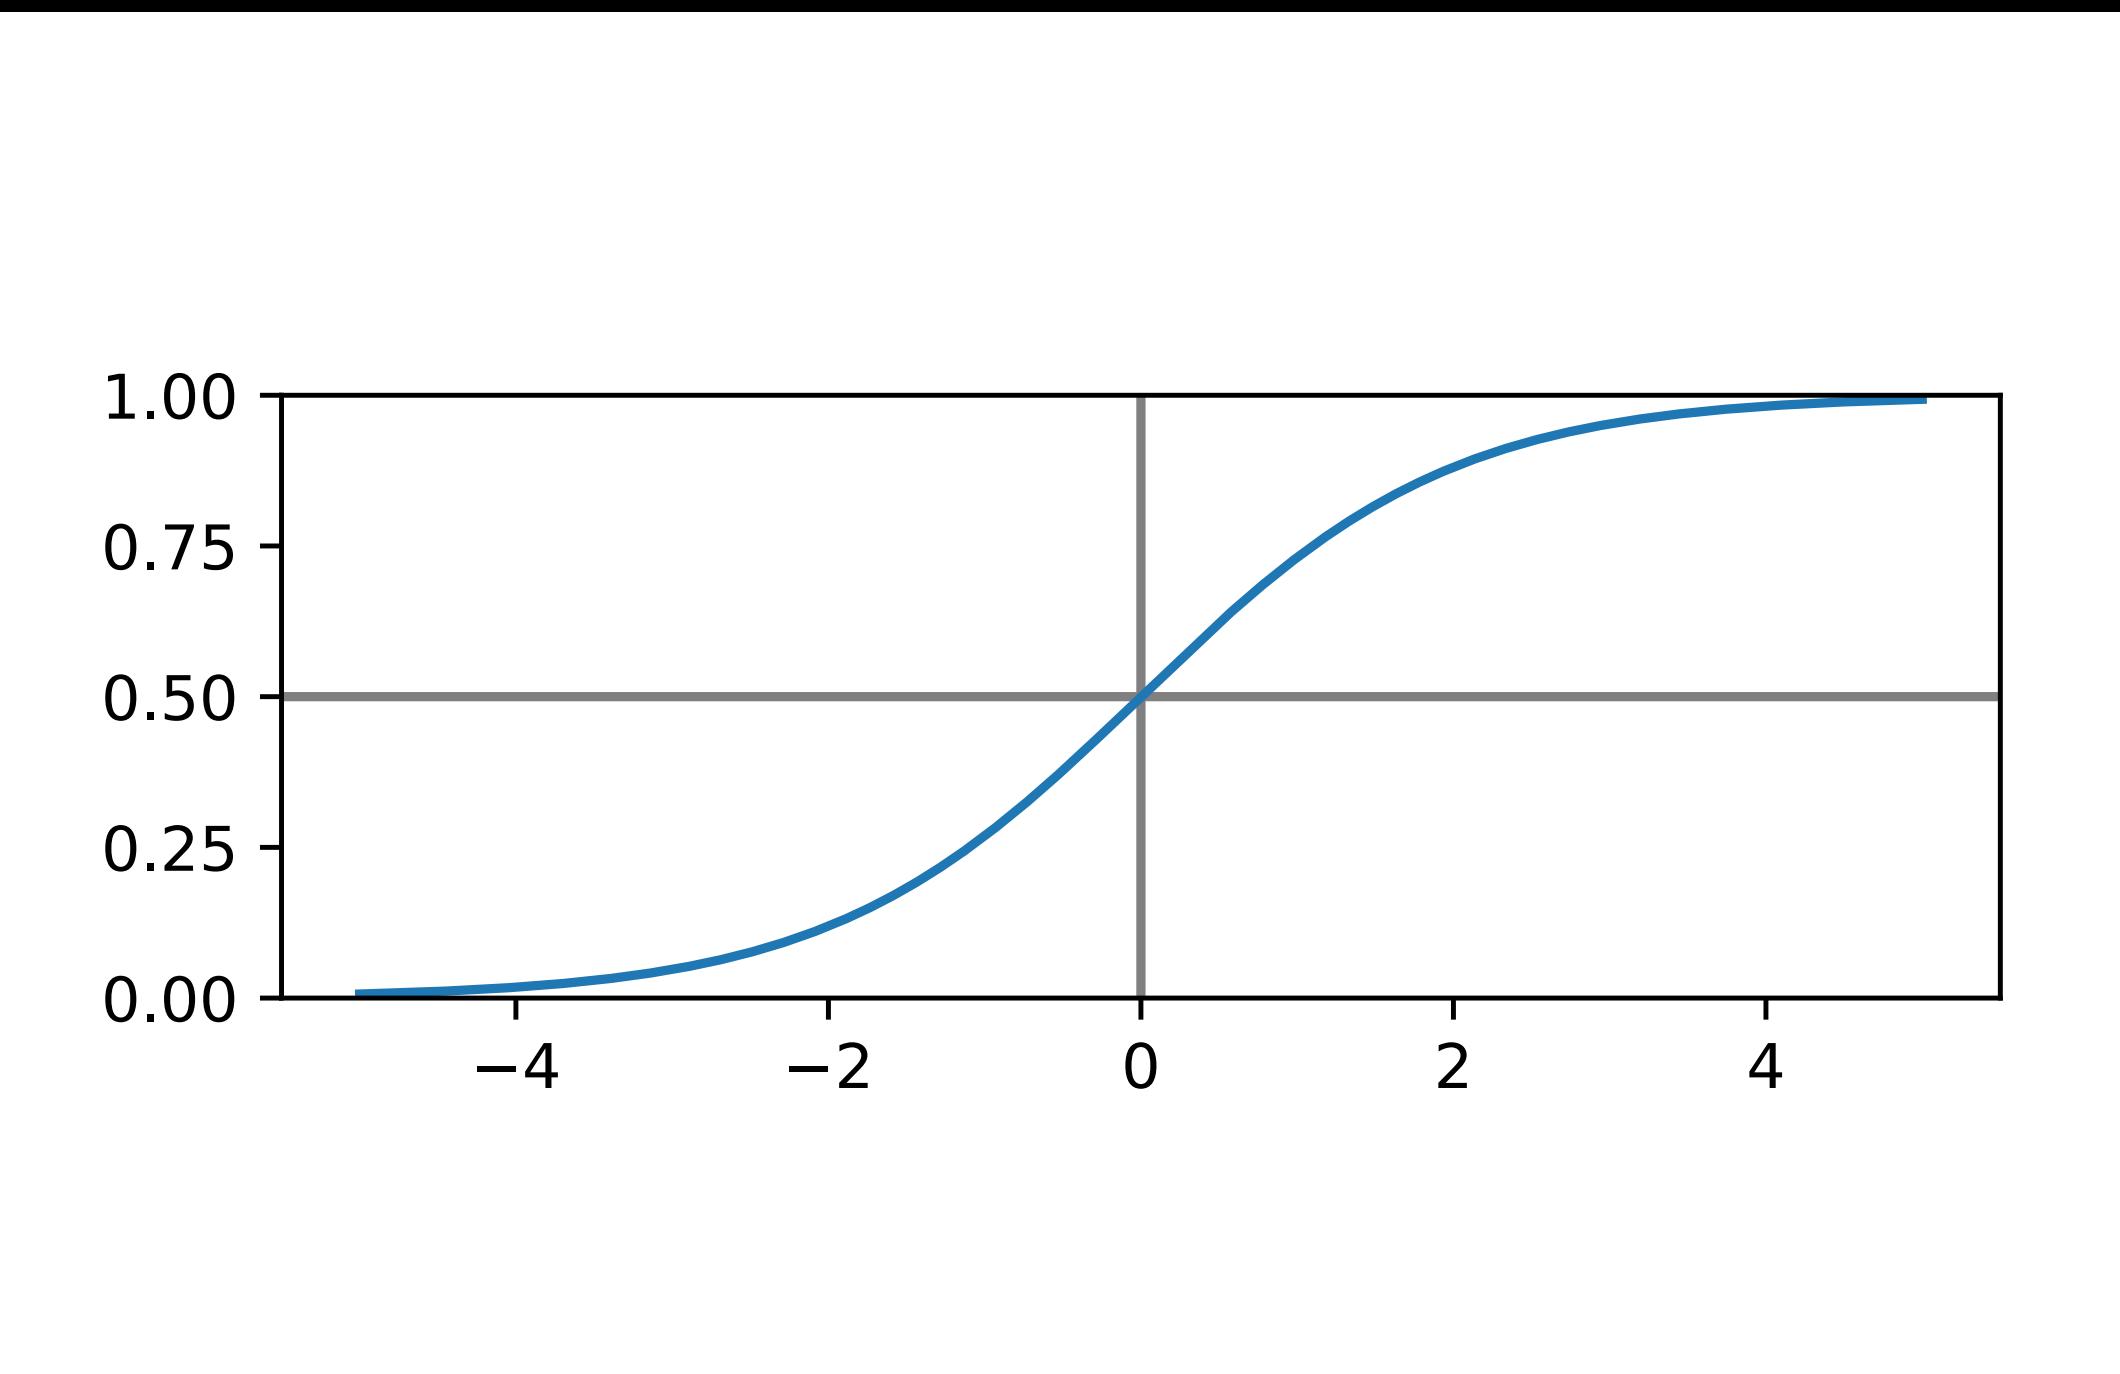
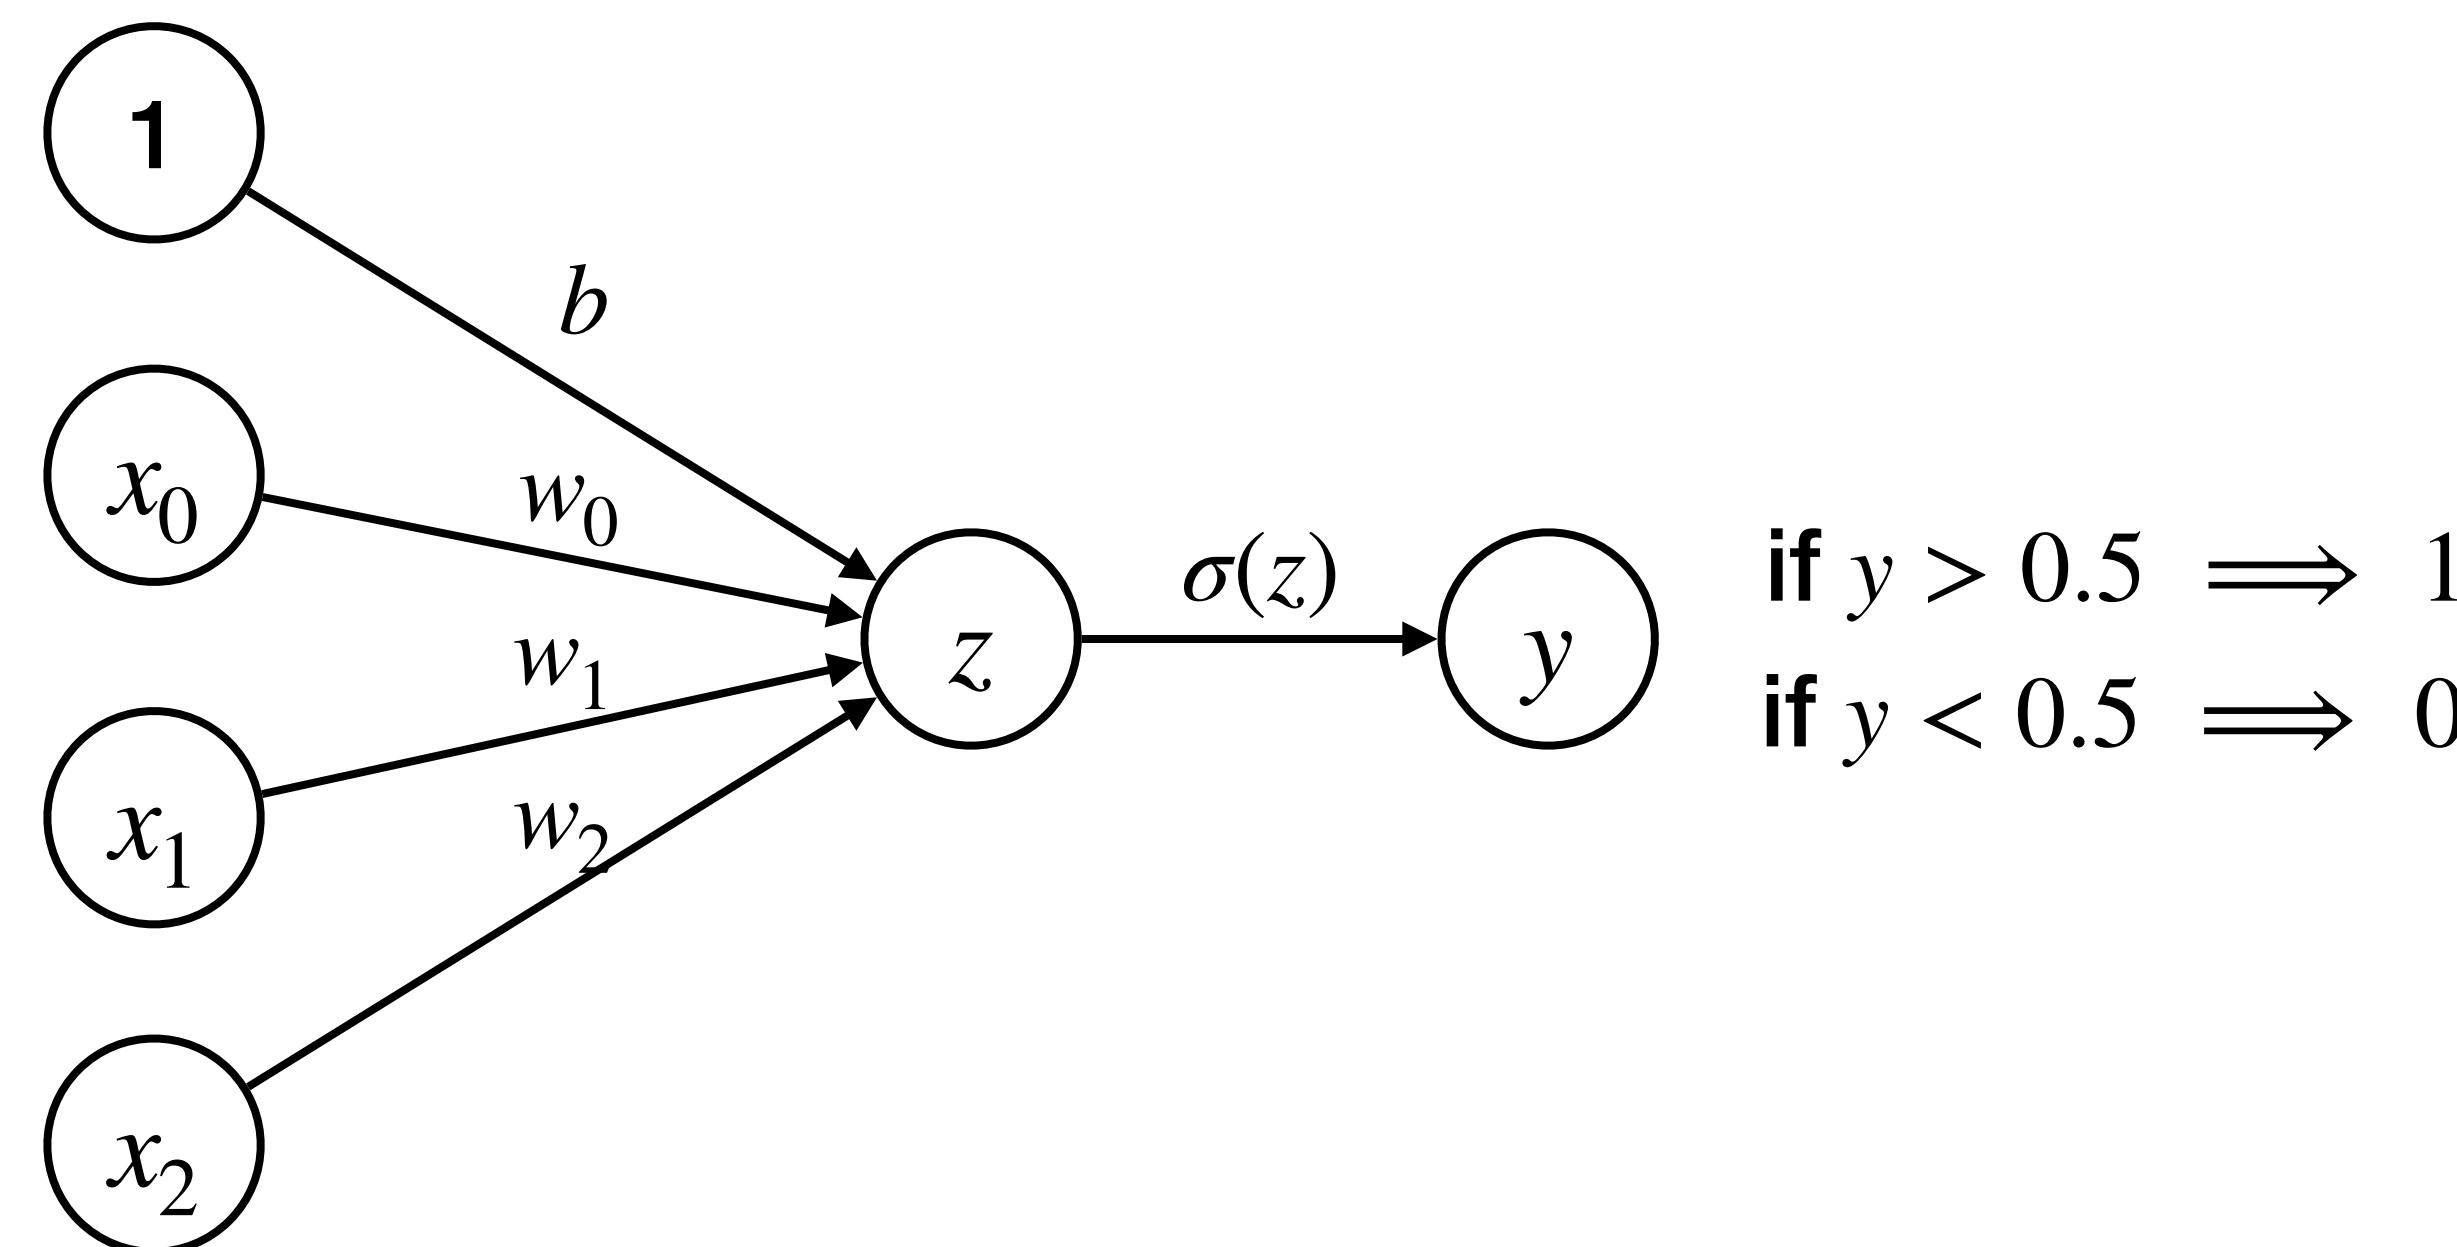
- Can we construct a mapping of $(-\infty, +\infty) \rightarrow \{0,1\}$?

Sigmoid function

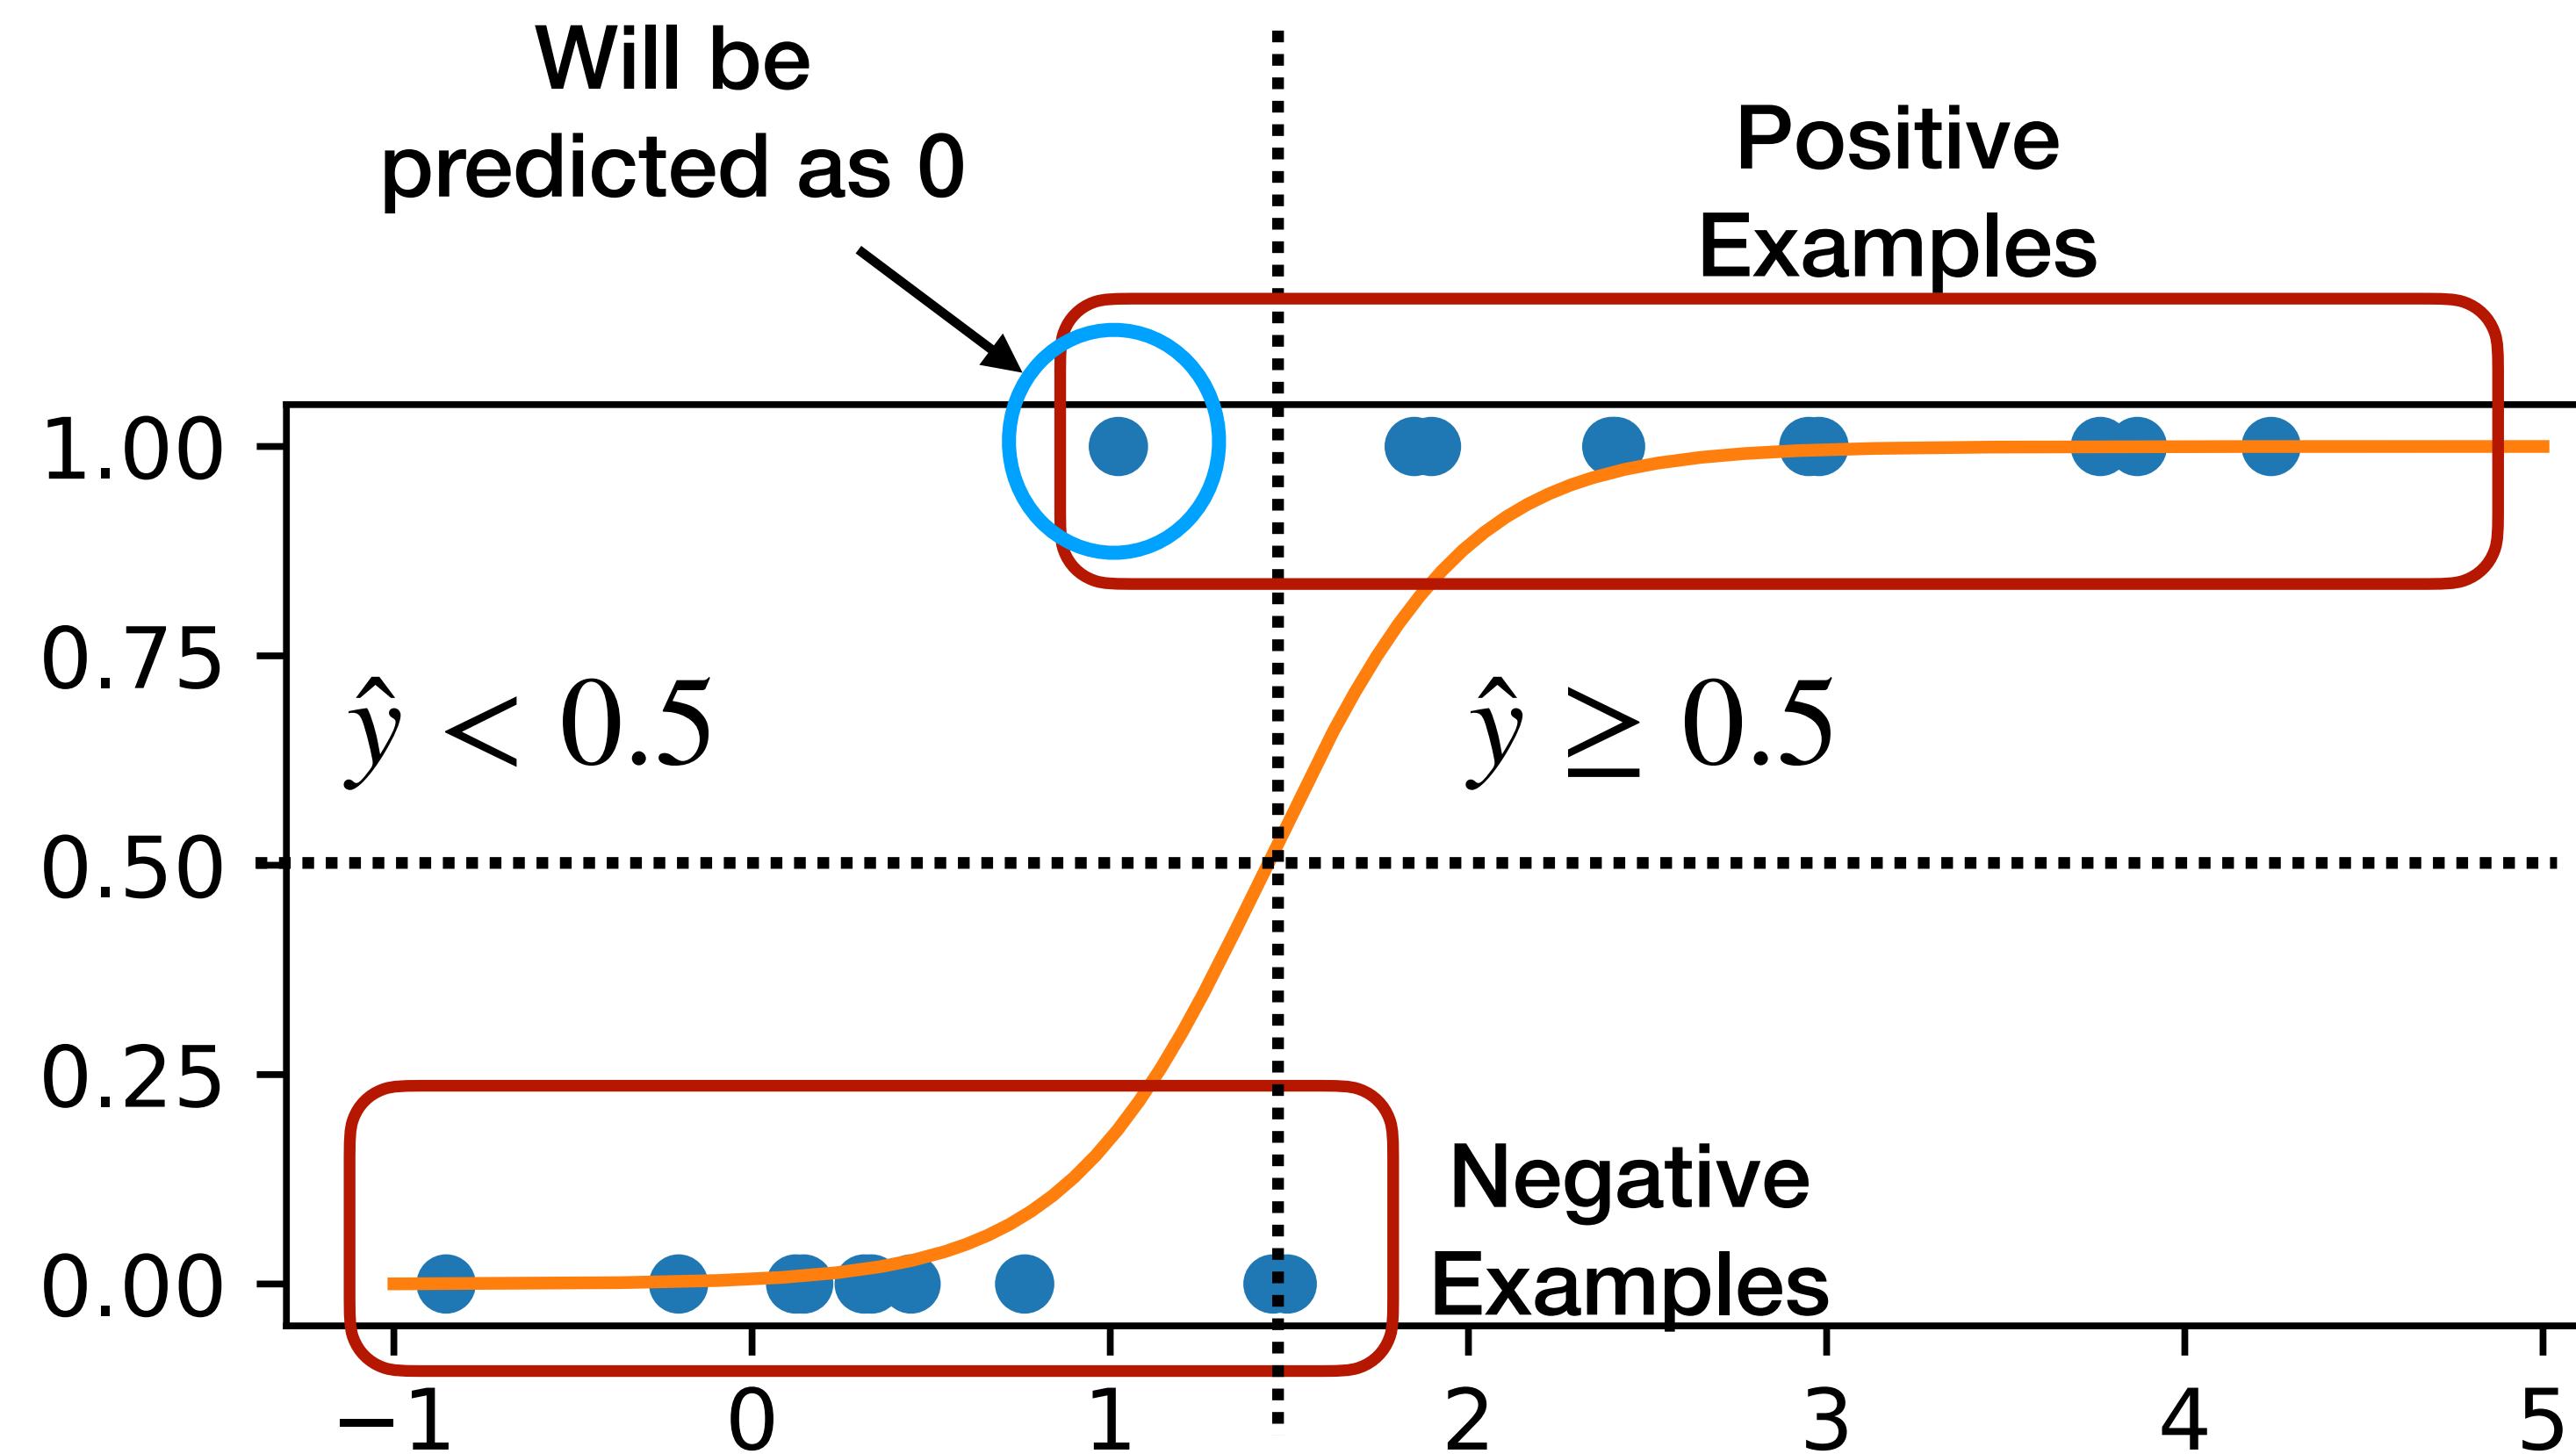
$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



“Transform” Linear Regression for Classification

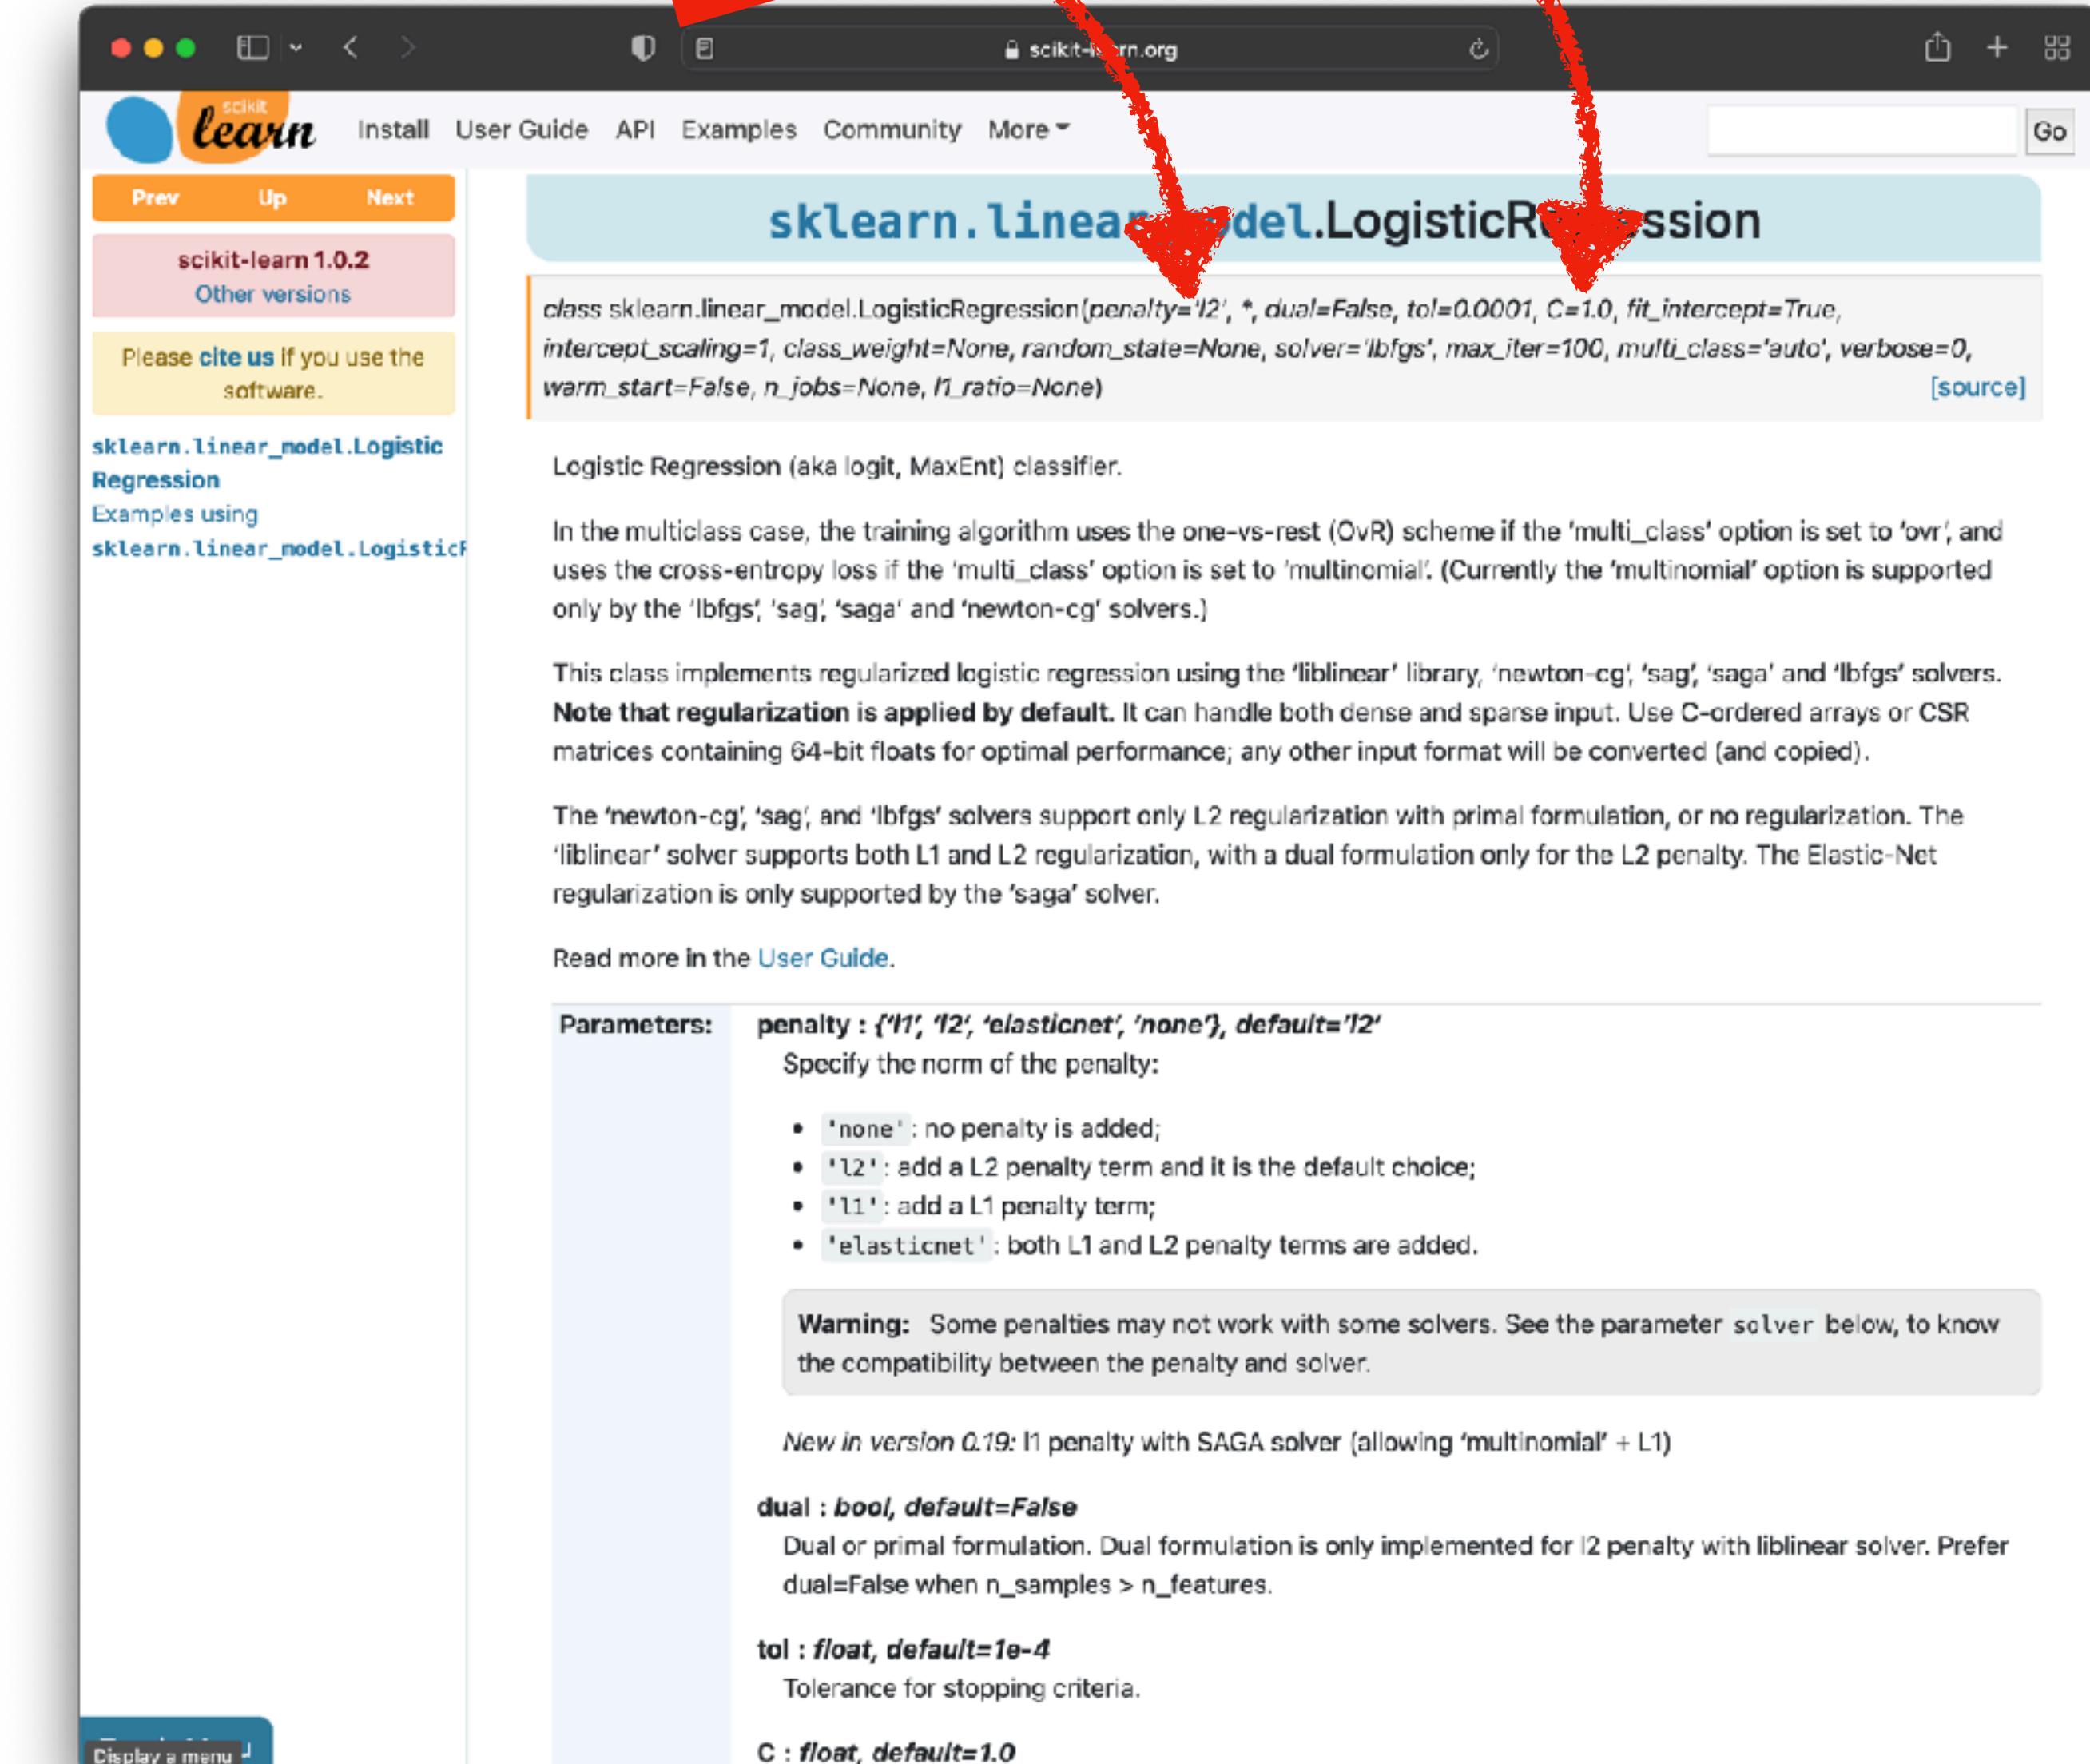


Logistic Regression: 1-D Example



Logistic Regression: Regularization

- L2 regularization is “on” by default in scikit-learn
 - Parameter C controls the strength of the regularization.



This Module's Learning Objectives

Part 2

Define regularization and its use of improving generalizability

Use linear regression models to predict scores for data elements

Explain how linear regression can be adapted for classification

Extract the class-label probabilities for model outputs

This Module's Learning Objectives

Part 2

Define regularization and its use of improving generalizability

Use linear regression models to predict scores for data elements

Explain how linear regression can be adapted for classification

Extract the class-label probabilities for model outputs

If Logistic Regression classifies labels via probability thresholds...

Can we extract these probabilities directly?

Indeed we can

`predict_log_proba(X)` [source]

Select logarithm of probability estimates.

The returned estimates for all classes are ordered by the label of classes.

Parameters: `X : array-like of shape (n_samples, n_features)`
Vector to be scored, where `n_samples` is the number of samples and `n_features` is the number of features.

Returns: `T : array-like of shape (n_samples, n_classes)`
Returns the log-probability of the sample for each class in the model, where classes are ordered as they are in `self.classes_`.

`predict_proba(X)` [source]

Probability estimates.

The returned estimates for all classes are ordered by the label of classes.

For a multi_class problem, if `multi_class` is set to be "multinomial" the softmax function is used to find the predicted probability of each class. Else use a one-vs-rest approach, i.e calculate the probability of each class assuming it to be positive using the logistic function. and normalize these values across all the classes.

Parameters: `X : array-like of shape (n_samples, n_features)`
Vector to be scored, where `n_samples` is the number of samples and `n_features` is the number of features.

Returns: `T : array-like of shape (n_samples, n_classes)`
Returns the probability of the sample for each class in the model, where classes are ordered as they are in `self.classes_`.

`score(X, y, sample_weight=None)` [source]

Return the mean accuracy on the given test data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

Parameters: `X : array-like of shape (n_samples, n_features)`
Test samples.

`y : array-like of shape (n_samples,) or (n_samples, n_outputs)`

Why would you want these probabilities?

Perhaps you only care about instances with high confidence

Instances with low confidence (prob. ~ 0.5) are interesting/hard samples

How else might we calculate these probabilities?

Train many slightly different models and assess predictions

Premise for “random forest” classification method

This Module's Learning Objectives

Part 2

Define regularization and its use of improving generalizability

Use linear regression models to predict scores for data elements

Explain how linear regression can be adapted for classification

Extract the class-label probabilities for model outputs

Questions?

Prof. Cody Buntain | @codybuntain | cbuntain@umd.edu
Director, Information Ecosystems Lab